

# Demystifying cloud-based data stores with Linux and Java

## **SESSION INTRO AND WHAT TO EXPECT**

Harold Wong  
Dev Experience – Azure Cloud

Bruno Terkaly  
Dev Experience – Azure Cloud

O'REILLY®

# Session Intro

## Session Flow

- A complete introduction to what Azure offers
- A specific emphasis on Linux-based workloads
- A deep introduction into automated provisioning of cloud resources in Azure
- Deploying Java Apps to the cloud
- Implementing Model-View-Controller Architectures in a web site with Azure Storage
- Working with Eclipse, Git, FileZilla, Maven, Azure SDKs

## Coverage of RDBMS systems

- SQL Server
- Azure SQL DB
- MySQL
- PostGres



## Coverage of NoSQL Data Stores

- Apache Cassandra
- MongoDB
- Redis
- DocumentDB



# Session Intro

<b>Activity</b>	<b>Start</b>	<b>Stop</b>	<b>Length</b>
Session Definition	1:30 PM	1:40 PM	10
What is Azure, Getting Ready, Installation	1:40 PM	1:50 PM	10
MySQL - Discuss	1:50 PM	2:02 PM	12
MySQL - Hands-On	2:02 PM	2:37 PM	35
PostGres - Discuss	2:37 PM	2:49 PM	12
PostGres - Hands-On	2:49 PM	3:24 PM	35
MongoDB - Discuss	3:24 PM	3:36 PM	12
MongoDB - Hands-On	3:36 PM	4:11 PM	35
Cassandra - Discuss & Wrap-up	4:11 PM	4:23 PM	12
Cassandra - Hands-On	4:23 PM	5:00 PM	37

# Session Intro

<b>Getting An Azure Subscription And Configuring Your Environment</b>	<b>01m 09s</b>
Provisioning An Ubuntu Linux VM At The Azure Portal	08m 37s
Installing The Azure Cross Platform Tooling	03m 29s
Provisioning Resources With The Azure Resource Manager - Part 1	06m 52s
Provisioning Resources With The Azure Resource Manager - Part 2	04m 57s
Provisioning Resources With The Azure Resource Manager - Part 3	14m 04s
Using The Azure Cross Platform CLI To Provision Infrastructure	07m 39s
Downloading And Installing Java, Eclipse, FileZilla And Git	11m 31s
Hello World With A Java Application And A Java Dynamic Web Application	12m 41s
Azure Storage - Introduction And Azure Blobs Overview - The Storage Landscape In Azure	06m 32s
Provisioning A Storage Account With The Portal And ARM	14m 53s

And much more online

# What do you want to cover?

Rank			DBMS	Database Model	Score		
May 2016	Apr 2016	May 2015			May 2016	Apr 2016	May 2015
1.	1.	1.	Oracle	Relational DBMS	1462.02	-5.51	+19.93
2.	2.	2.	MySQL 	Relational DBMS	1371.83	+1.72	+77.56
3.	3.	3.	Microsoft SQL Server	Relational DBMS	1142.82	+7.77	+11.79
4.	4.	4.	MongoDB 	Document store	320.22	+7.78	+42.90
5.	5.	5.	PostgreSQL	Relational DBMS	307.61	+3.89	+34.09
6.	6.	6.	DB2	Relational DBMS	185.96	+1.87	-15.09
7.	↑ 8.	↑ 8.	Cassandra 	Wide column store	134.50	+4.83	+27.95
8.	↓ 7.	↓ 7.	Microsoft Access	Relational DBMS	131.58	-0.39	-14.00
9.	9.	↑ 10.	Redis 	Key-value store	108.24	-3.00	+13.51
10.	10.	↓ 9.	SQLite	Relational DBMS	107.26	-0.70	+2.10

# Tools that we will download and install

<b>Java</b>	<ul style="list-style-type: none"><li>• An SDK and Runtime to support Java Programming</li></ul>
<b>Eclipse</b>	<ul style="list-style-type: none"><li>• Developer tooling we will use to write, debug, and deploy Java apps</li></ul>
<b>FileZilla</b>	<ul style="list-style-type: none"><li>• FTP client supported by Windows, MacOS, and Linux.</li><li>• Used for deployment.</li></ul>
<b>Git</b>	<ul style="list-style-type: none"><li>• Used for source code control and for deployment</li></ul>
<b>Azure SDKs</b>	<ul style="list-style-type: none"><li>• Various SDKs will be used for specific tasks in Azure</li></ul>
<b>Jersey – Java API</b>	<ul style="list-style-type: none"><li>• Java libraries useful for REST-based programming</li></ul>
<b>Tomcat</b>	<ul style="list-style-type: none"><li>• Apache web server used to run our web applications. Also hosted in Azure.</li></ul>

# Compiling Maven Code

3 main commands with Eclipse to compile project:

- Right mouse click on the project and select:
  - Maven / Update Project
  - Run As / Maven Clean
  - Run As / Maven Install

The 3 crucial commands in Eclipse when working with Maven

# Getting to the downloads

<http://bit.ly/azuresessiondownloads>  
<http://bit.ly/osconcode>

## Download Notes

You will need two versions of Java.

- The latest 8.0 version for running Eclipse.
- You will need 1.7.0\_51 for our Java Web Apps because that is the version hosted in Tomcat.

## 64-bit

You will need to make sure you download the 64-bit versions of everything if you are running a 64-bit Operating system.

### Java

<http://www.oracle.com/technetwork/java/javase/downloads/jdk8-downloads-2133151.html>

<http://www.oracle.com/technetwork/java/javase/downloads/java-archive-downloads-javase7-521261.html>

<http://www.oracle.com/technetwork/java/javase/downloads/java-archive-downloads-javase7-521261.html#jre-7u51-oth-JPR>

### Eclipse

<http://www.eclipse.org/downloads/packages/eclipse-ide-java-ee-developers/marsr>



HowToCreateAVM.mp4

## FileZilla

<https://filezilla-project.org/download.php>

## Git

<http://git-scm.com/download>

## Azure SDKs

Download as needed throughout the course

## Jersey – Java API

<https://jax-rs-spec.java.net/>

## Tomcat

<http://tomcat.apache.org/download-70.cgi>

# Opening up ports on your Azure VM

Inbound security rules

Add Default rules

PRIORITY	NAME	SOURCE	DESTINATION	SERVICE	ACTION	...
1000	default-allow-ssh	Any	Any	SSH (TCP/22)	Allow	...
1010	mysqlport	Any	Any	MySQL (TCP/3306)	Allow	...
1020	postgres	Any	Any	Custom (ANY/5432)	Allow	...
1030	mongodb	Any	Any	Custom (ANY/27017)	Allow	...
1040	cassandra1	Any	Any	Custom (ANY/7000)	Allow	...
1050	cassandra2	Any	Any	Custom (ANY/9160)	Allow	...
1060	cassandra3	Any	Any	Custom (ANY/9042)	Allow	...
1070	cassandra4	Any	Any	Custom (ANY/7199)	Allow	...

# Session Intro

What this session is focused on

- This session is focused on the world's most popular data stores using the most popular programming language

This session provides high level plus running Java code

- This session is all about concrete examples, both from the perspective of installation and configuration, as well as creating code that directly manipulates data



Windows Azure  
Storage



Storage BLOB



Storage Table

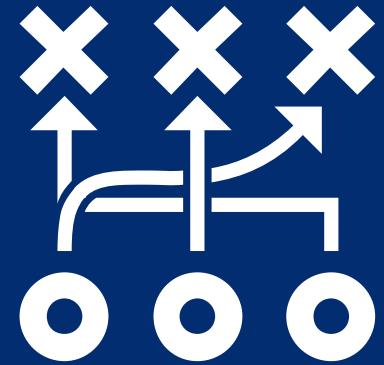


Storage  
Queue

# Session Intro

## Session Philosophy

- No copy and paste
- Everything starts with File / New Project
- Code is typed in one character at a time
- Fastest, most productive way to learn
- Full Source code provided



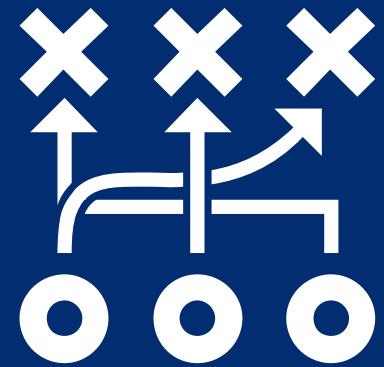
## Learning Strategy

- Students follow along from beginning to end
- Encapsulates all 3 learning techniques
  - Hearing
  - Watching
  - Doing

# Session Intro

## Session Philosophy

- No copy and paste
- Everything starts with File / New Project
- Code is typed in one character at a time
- Fastest, most productive way to learn
- Full Source code provided

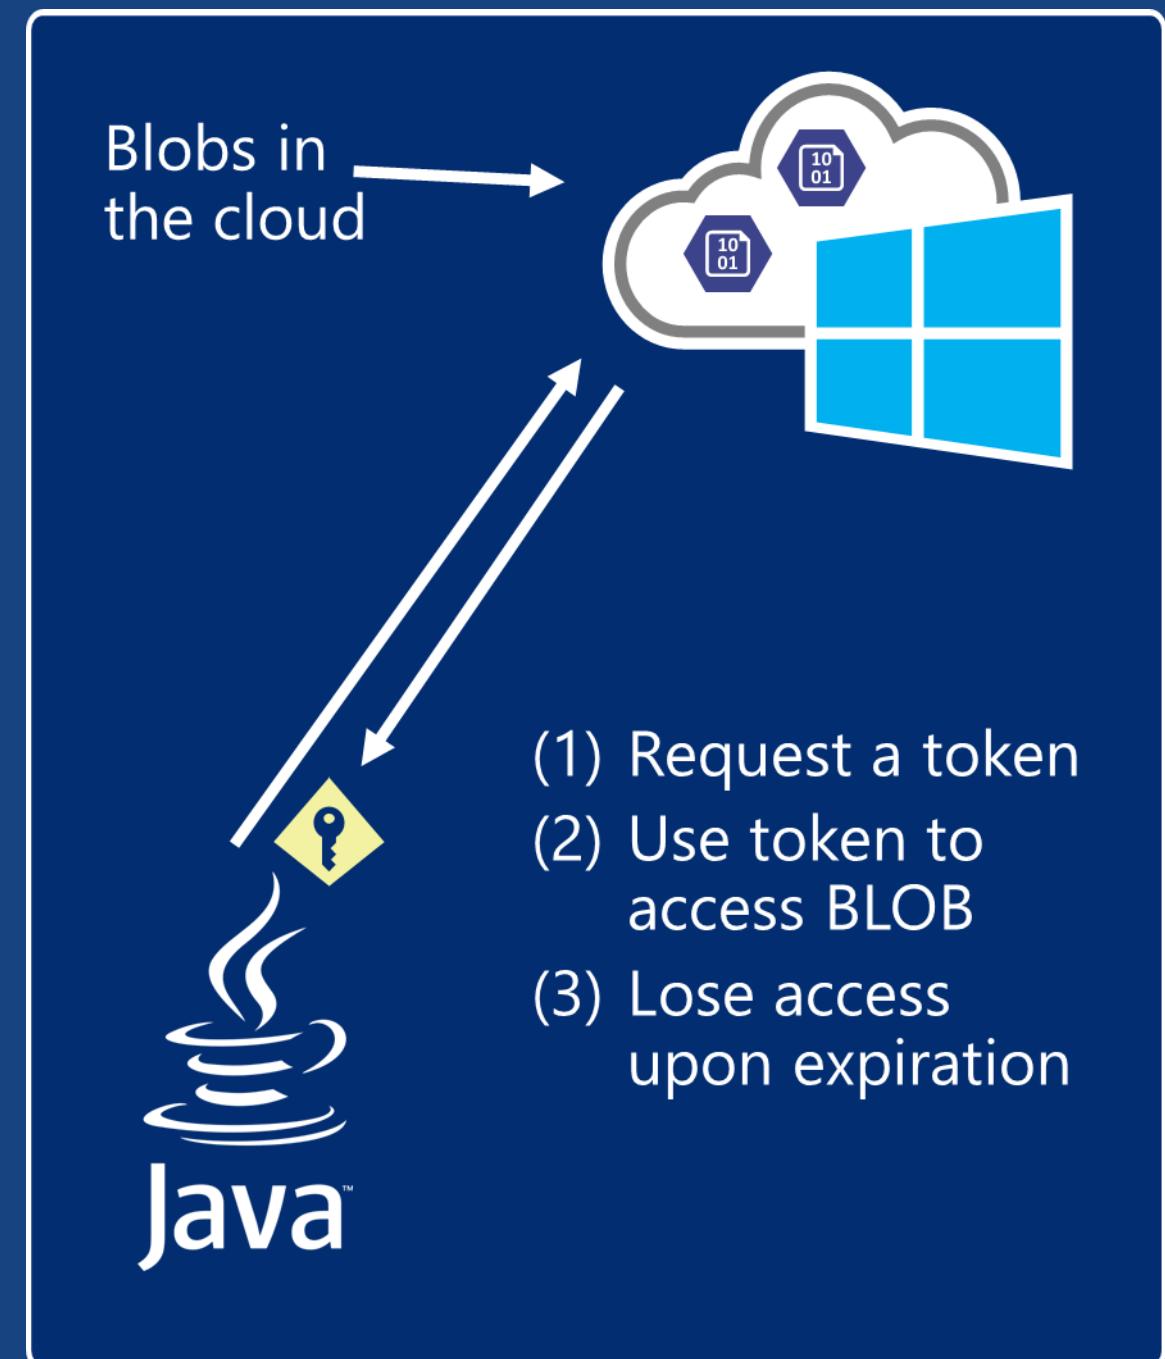
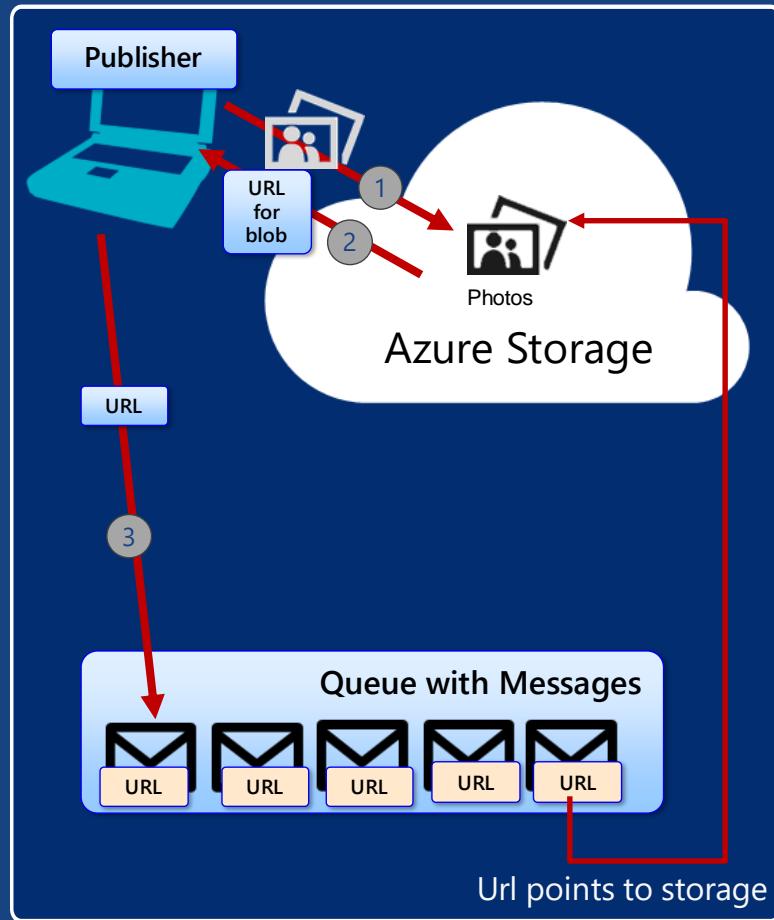


## Learning Strategy

- Students follow along from beginning to end
- Encapsulates all 3 learning techniques
  - Hearing
  - Watching
  - Doing

# Session Intro

For visual learners, numerous concise, clear diagrams are included



# Session Intro

This session is hands-on

- Concrete examples for both installation and coding
- Demos on live, running Linux VMs
- All commands reviewed and explained
- Step by step Java samples using Eclipse, including debugging



Cloud enabled examples



Hands-on construction

- Great for anybody getting started and cloud computing should consider this session
- The focus is to be practical and for students to start writing meaningful code immediately
- Great for anybody looking to land a new job in the cloud

# Bruno Terkaly

- Principal Software Engineer at Microsoft with a 10 Year tenure
- A long career in education, writing, speech giving
- Popular and content-heavy blog at <http://blogs.msdn.com/b/brunoterkaly/>
- Developer and publisher of multiple phone apps in the Windows Store
- Finalist in InfoWorld's best 12 youtube.com instructional videos in 2013
  - <http://www.infoworld.com/article/2606686/it-training/114051-From-IFTTT-to-Node.js-A-bakers-dozen-of-tech-tutorials.html>

The screenshot shows a web browser window with the URL <https://msdn.microsoft.com/en-us/magazine/mt149362>. The page is titled "Bruno Terkaly". It displays three articles by Bruno Terkaly:

- Azure Insider - Creating Unified, Heroku-Style Workflows Across Cloud Platforms** (Sep 2015) - A summary of the third installment in the Azure Insider series, which explores how technologies like DocumentDB and Node.js can be used to create unified workflows across different cloud platforms.
- Event Hubs for Analytics and Visualization, Part 3** (Jul 2015) - A summary of the third part of a series on Event Hubs, focusing on using Event Hubs to store data and visualize it on a Windows Phone.
- Event Hubs for Analytics and Visualization, Part 2** (Jun 2015) - A summary of the second part of the series, which discusses the challenges of dealing with massive data stores and how Event Hubs can help.

Each article has a "Read article" button below it.

Mobile, Node.js, Big Data, Multi-Tenant Architectures, Azure, Windows Phone, Open Source, C/C++, C#, Node.js/Javascript, SQL, LINQ, and more

# Bruno Terkaly

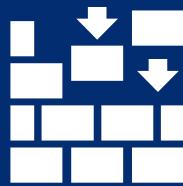
- Co-organizer of a 6,000+ member strong Bay Area Software Engineers on meetup.com
- Create huge live events and speaking to 1,000's of developers in the Silicon Valley Area
- Topics include Robotics, Machine Learning, Data Science, Node.js, Cloud Computing and much more (see <http://www.meetup.com/software>)
- Assist and directly engage with some of the largest software brands in the world to help them create products that run on Azure
  - Docker
  - Github
  - Mesosphere
  - DataStax
  - CoreOS
  - Deis



Top scoring presenter at the Microsoft Executive Briefing Center in Redmond, and Silicon Valley Code Camp



Work closely with Microsoft/Redmond-based engineering teams to accelerate customer adoption of Azure



Influenced roadmap for future versions of Azure



Organize and participated in weeklong engineering migration efforts for high-value customers.

# AZURE FEATURE OVERVIEW – PART 1

BRUNO TERKALY – PRINCIPAL SOFTWARE ENGINEER – MICROSOFT - AZURE

O'REILLY®

# Introduction to Azure

- A public cloud computing platform
- For building, deploying, and managing applications and services
- Supports extensive amount of languages, frameworks, tools



Public Cloud



Global Set of Data Centers



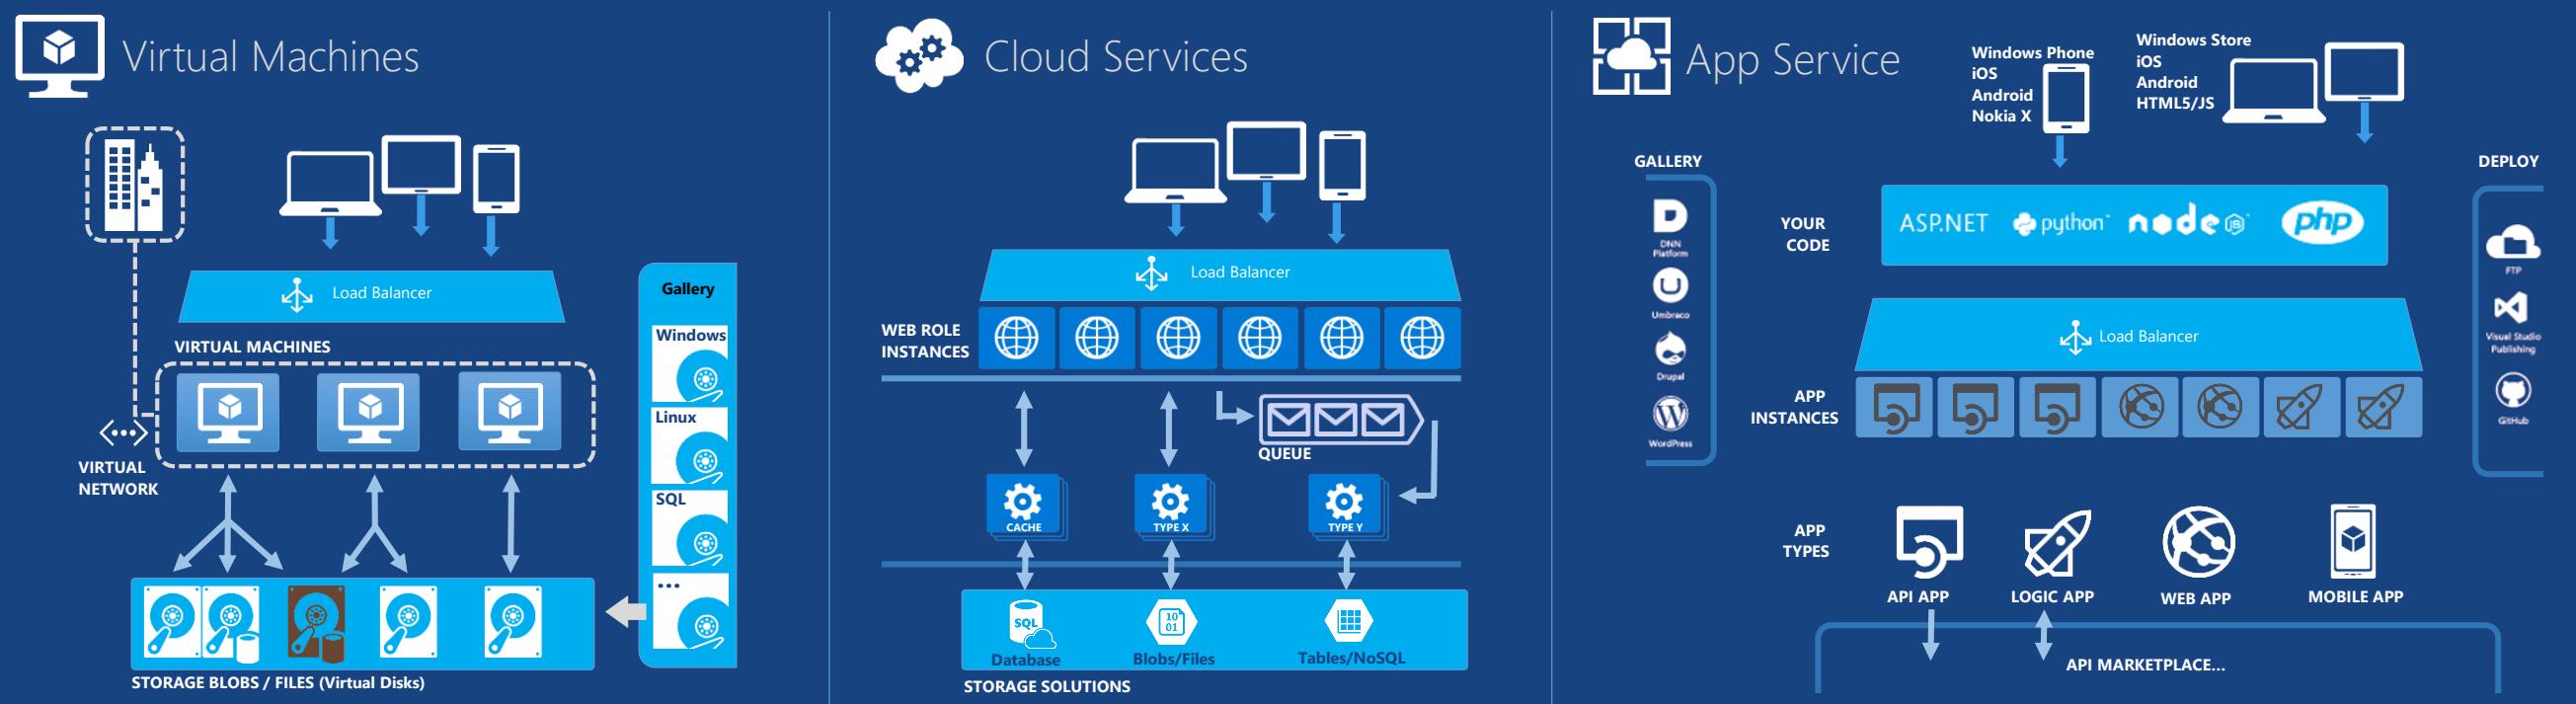
Provides an extensive compute infrastructure



Web Sites/Services



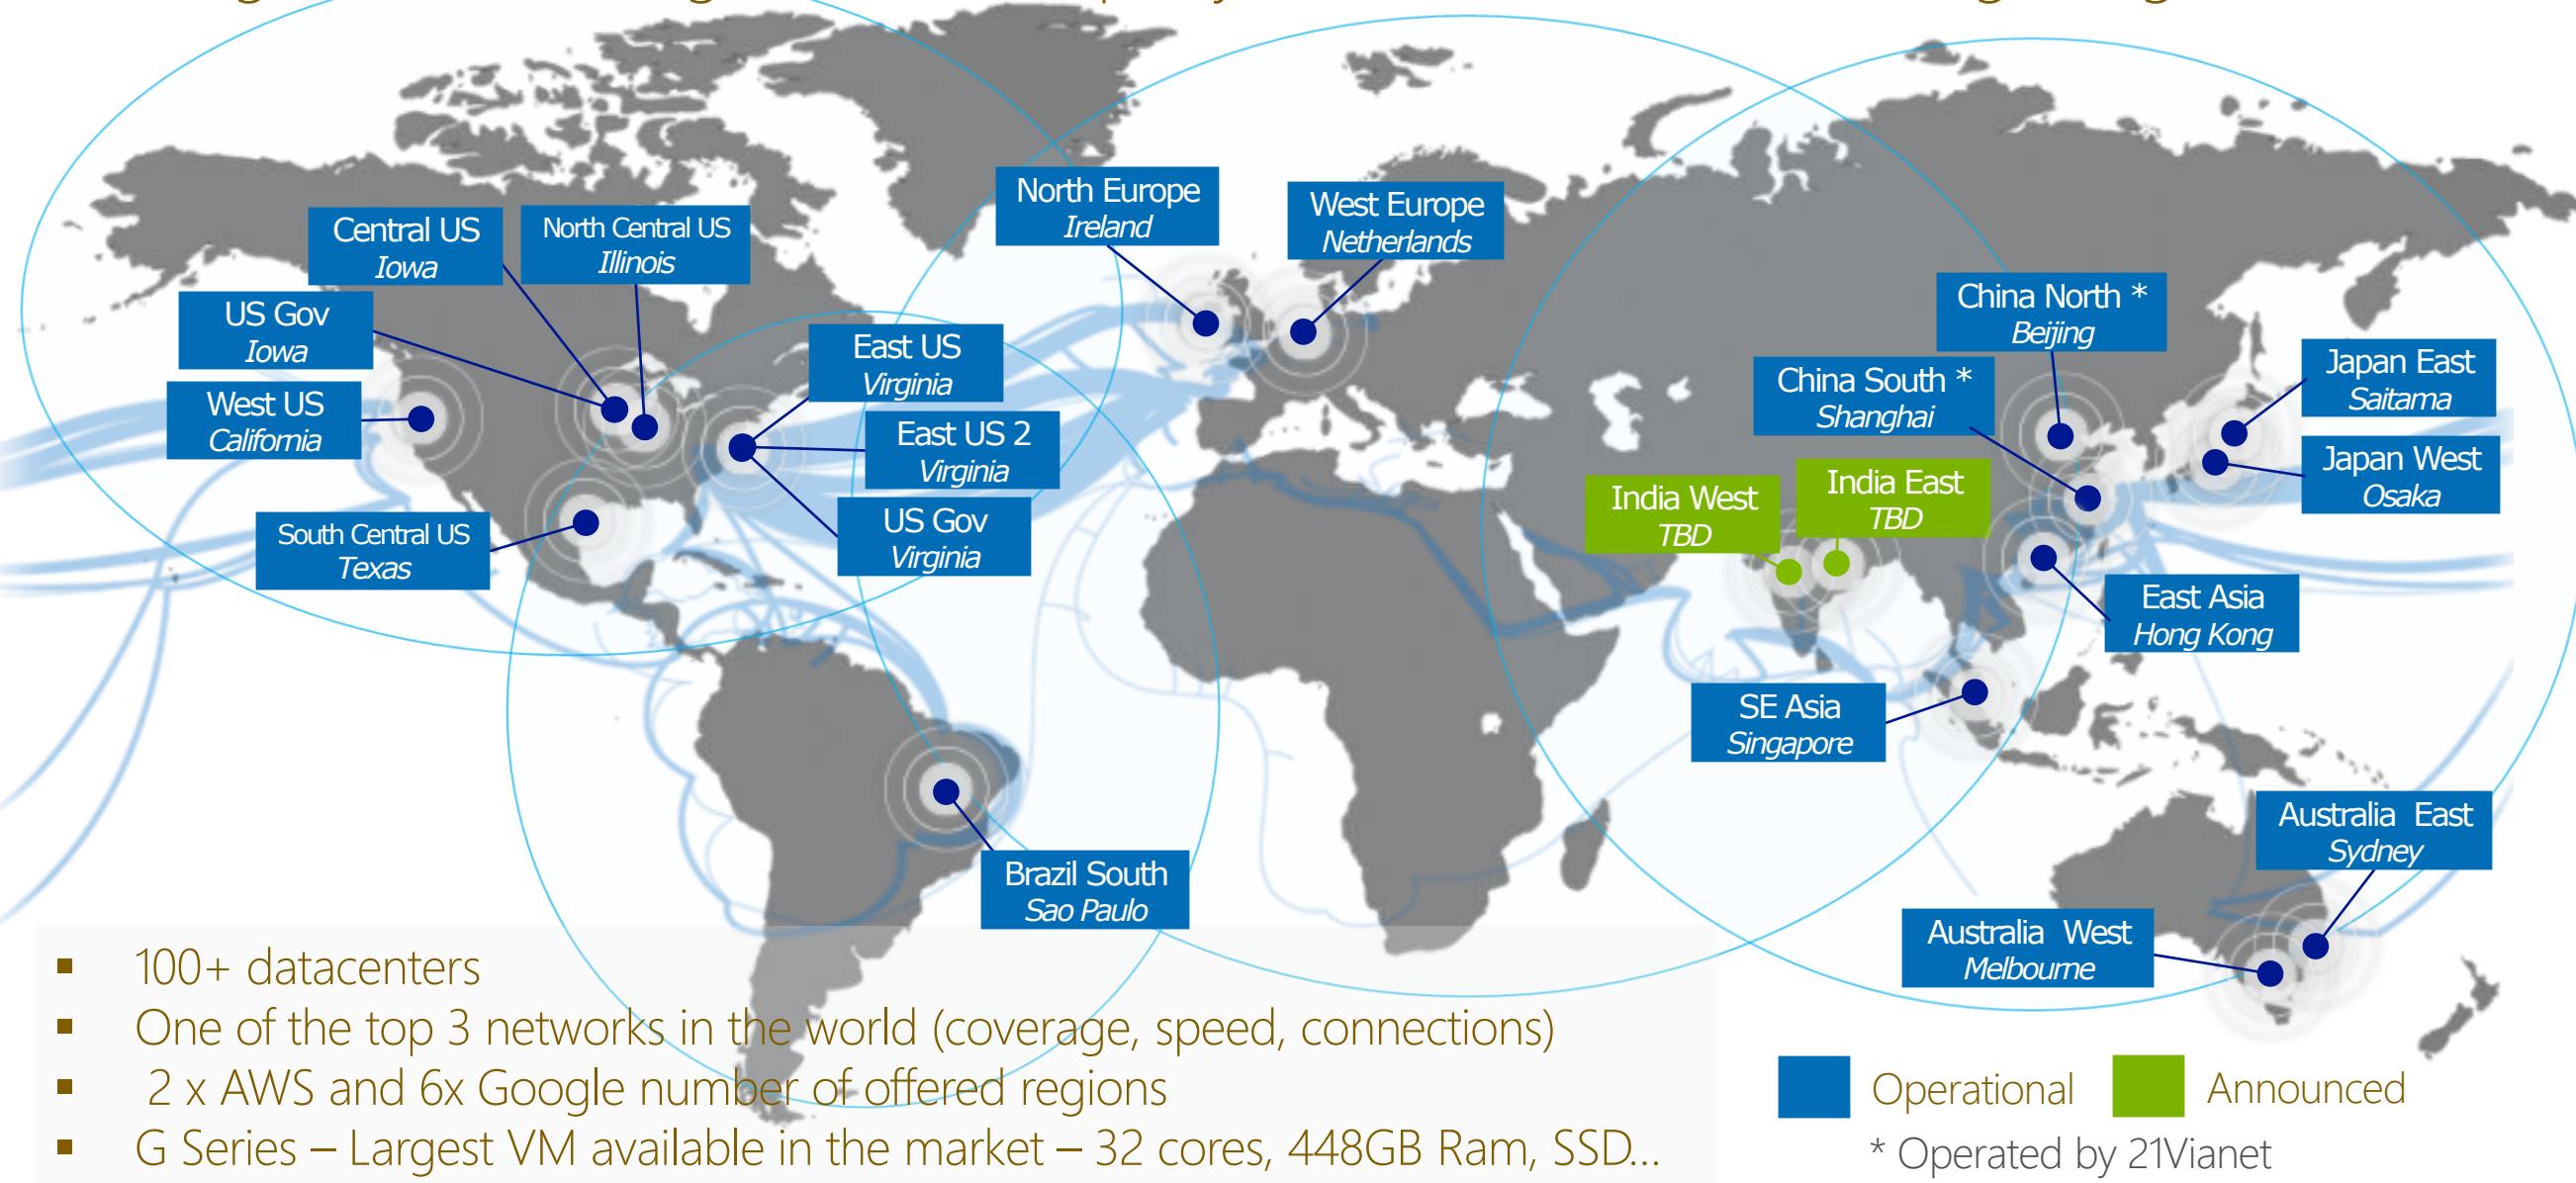
Data



COMPUTE	NETWORKING	IDENTITY & ACCESS	MEDIA & CDN
<b>Virtual Machines</b> Get full control over a server in the cloud and maintain it as your business requires.	<b>Cloud Services</b> Managed Virtual Machines with specific web and worker roles that are stateless.	<b>Batch</b> For running large scale parallel and high performance computing (HPC) applications.	<b>Scheduler</b> Create jobs that run reliably on simple or complex schedules to invoke any type of service.
<b>Web &amp; Mobile</b> <b>Web Apps</b> Managed web platform, get started for free and scale as you go using many tools/languages.	<b>Mobile Apps</b> Add backend capabilities to mobile apps, with native client support on most device platforms.	<b>API Apps</b> Create and surface your app logic as APIs for other services and apps to consume.	<b>Logic Apps</b> Build/execute business processes by linking your own custom API's with an API Gallery/Marketplace
<b>Storage &amp; Backup</b> <b>Storage Blobs &amp; Files</b> Store binary application data and web content – store for dedicated and shared virtual disks for VM's	<b>Backup</b> Managed service that handles backup/restore of Windows Server machines/backup agent.	<b>Import/Export</b> For massive data transfer – ship encrypted disks to move data in/out of blob storage.	<b>Site Recovery</b> Coordinate replication and recovery of System Center private clouds
<b>Hybrid Integration</b> <b>Storage Queues</b> Simple message queue for application de-coupling architecture for scale out.	<b>Biztalk Services</b> Build EDI and Enterprise App Integration (EA) solutions in the cloud.	<b>Hybrid Connections</b> Connect apps in Azure with on-premises resources without a VPN or dedicated line.	<b>Service Bus</b> Messaging capabilities (pub/sub, queues) and on-premises to cloud connectivity solution.
DATA	ANALYTICS	DEVELOPER SERVICES	
<b>SQL Database</b> Managed relational database service with high availability and selectable performance levels.	<b>HDInsight</b> Big Data (based on Apache Hadoop) analytics that integrate easily with Microsoft Office.	<b>Machine Learning</b> Mine historical data with compute power to predict future trends or behavior.	<b>Media Services</b> Range of services that support video on-demand and live streaming workflows.
<b>DocumentDB</b> Store/retrieve millions of JSON objects from a highly scalable NoSQL document database.	<b>Stream Analytics</b> Process data streams in real-time to discover and react to trends.	<b>Data Factory</b> Ingest data from multiple sources to combine into a cloud based Data Warehouse.	<b>Content Delivery Network (CDN)</b> Cache content for your apps at 100's of edge locations to improve user experiences.
MANAGEMENT	PORTALS	COMMERCE	
<b>Automation</b> Run durable PowerShell scripts to automate frequent, long running, complex Azure tasks.	<b>Portal</b> Web based experience to provision, control and monitor all Azure services.	<b>Key Vault</b> Safeguard and control keys and secrets in cloud scale hardware security modules.	<b>Store / Marketplace</b> Find and manage other services provided by third parties.
<b>Hybrid Connections</b> Connect apps in Azure with on-premises resources without a VPN or dedicated line.	<b>Operational Insights</b> Analyze and troubleshoot on-premises IT infrastructure without using instrumented code.	<b>VM Depot</b> Find free open source VM images that you can download and run in Azure Virtual Machines.	<b>Application Insights</b> Analyze app usage, availability and performance to detect issues and solve problems proactively.

# Huge infrastructure scale is the enabler

19 Regions ONLINE...huge datacenter capacity around the world...and we're growing



# Analytics



## HDI insight

Big Data (based on Apache Hadoop) analytics that integrate easily with Microsoft Office.



## Machine Learning

Mine historical data with compute power to predict future trends or behavior.



## Stream Analytics

Process data streams in real-time to discover and react to trends.



## Data Factory

Ingest data from multiple sources to combine into a cloud based Data Warehouse.



## Event Hubs

Ingest, persist, process millions of events per second from millions of devices.



## Mobile Engagement

About understanding users and their behavior with mobile devices. Provides analytics on how users interact with mobile applications.

# Linux Focus

- Cross-Plat CLI scripting, not PowerShell
- Java, not C#
- MySQL and SQL Server
- Java SDKs
- Hadoop, Pig, Hive
- Services via REST

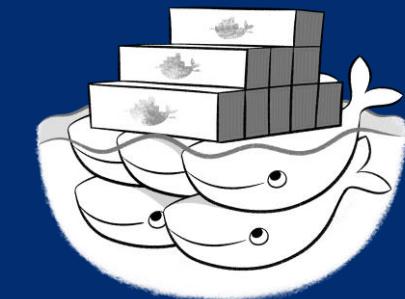
W3C



Open Standards



Hadoop



Pig  
Hive

Docker

# Getting Ready To Write Code

- We will:
  - Download software
  - Install Software
  - Configure Dev Environment

This is a demo heavy module

- 1 Java
- 2 Eclipse
- 3 FileZilla
- 4 Git
- 5 Azure SDKs
- 6 Tomcat
- 7 Jersey – Java API

# Tools that we will download and install

<b>Java</b>	<ul style="list-style-type: none"><li>• An SDK and Runtime to support Java Programming</li></ul>
<b>Eclipse</b>	<ul style="list-style-type: none"><li>• Developer tooling we will use to write, debug, and deploy Java apps</li></ul>
<b>FileZilla</b>	<ul style="list-style-type: none"><li>• FTP client supported by Windows, MacOS, and Linux.</li><li>• Used for deployment.</li></ul>
<b>Git</b>	<ul style="list-style-type: none"><li>• Used for source code control and for deployment</li></ul>
<b>Azure SDKs</b>	<ul style="list-style-type: none"><li>• Various SDKs will be used for specific tasks in Azure</li></ul>
<b>Jersey – Java API</b>	<ul style="list-style-type: none"><li>• Java libraries useful for REST-based programming</li></ul>
<b>Tomcat</b>	<ul style="list-style-type: none"><li>• Apache web server used to run our web applications. Also hosted in Azure.</li></ul>

# The SQL Spectrum

- Control versus Cost
- Pay-as-you-go
- Upfront investment
- Bring your own license
- Multi-tenancy
- Scaling, Backups, Upgrades, DR
- Time to market
- SLA (VMs = 99.95%, Azure SQL DB = 99.99%)
- Compatibility

Physical

Raw Iron

Virtual

SQL Server Private  
Cloud with Virtual  
Machines

IaaS

SQL Server in an  
Azure VM

PaaS

Azure SQL  
Database

# Overview of Relational Storage

## Provisioning Options

- Oracle
- SQL Server
- MySQL
- Azure SQL Database

**ORACLE®**

Oracle 12.1.0.1 Enterprise

Oracle Linux 6.4

Oracle Linux 7.0

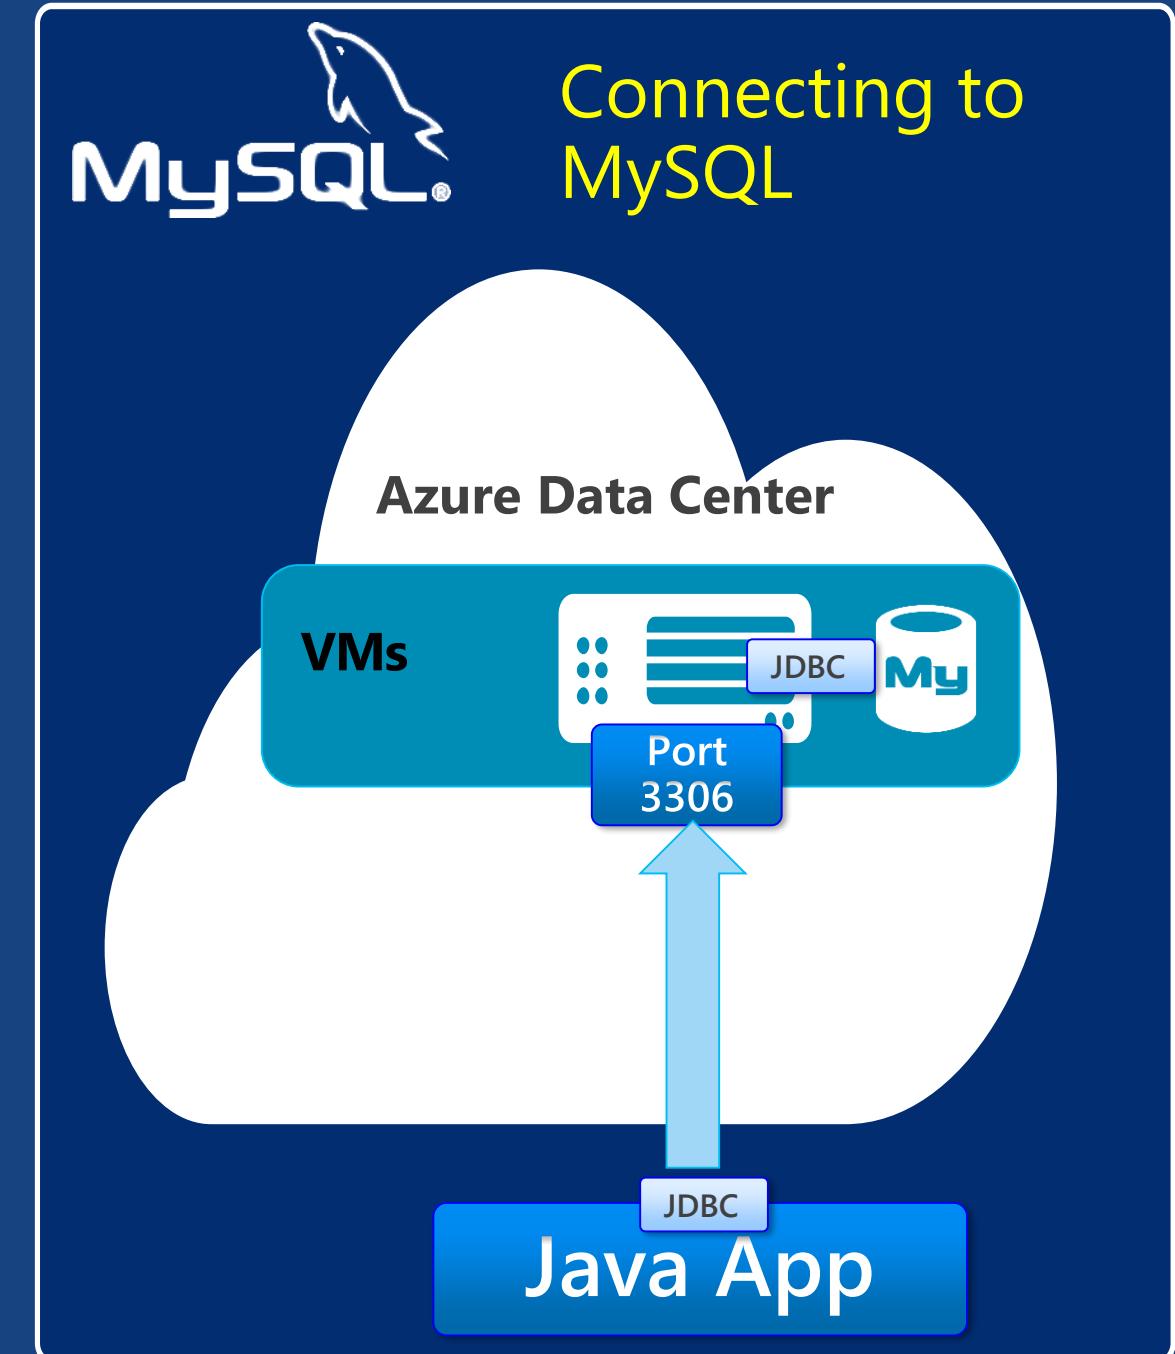
Oracle Web Logic

# **CONNECTING TO MYSQL FROM JAVA – PART 1**

**O'REILLY®**

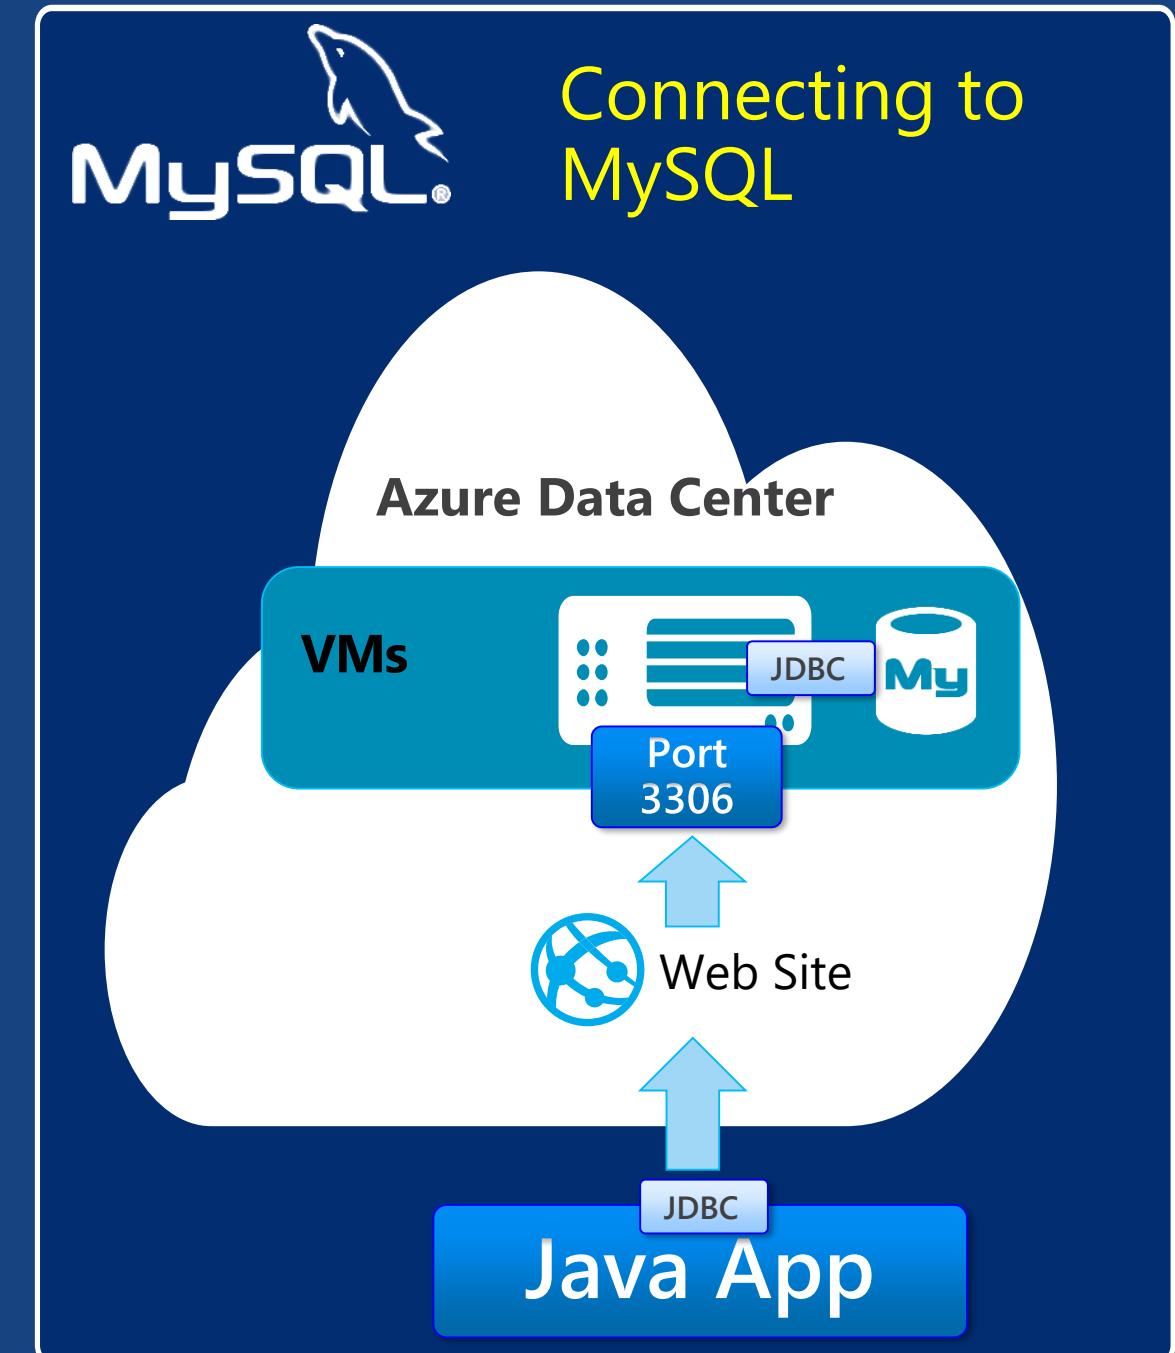
# Connecting to MySQL from Java

- MySQL listens to port 3306
- It isn't open by default
- We will use the Azure Cross Platform Tooling to open



# Connecting to MySQL from Java

- A better architecture
- Going through a web site



# Setting Up MySQL/JDBC Driver on Ubuntu

```
$ sudo apt-get install libmysql-java
```

JDBC drivers for Ubuntu

# Opening up port 3306

```
$ azure vm endpoint create oreillytestvm 3306 3306
```

Name of VM (you can see it on the command line)

Port that MySQL listens to

public port local port

# Opening a connection to the database

- A few modifications are needed
- Granting privileges
- Modifying my.cnf, changing to 0.0.0.0

```
# Get to a mysql prompt  
$  
GRANT ALL PRIVILEGES ON *.* TO  
'root'@'localhost' IDENTIFIED BY 'root' WITH  
GRANT OPTION;  
flush privileges;
```

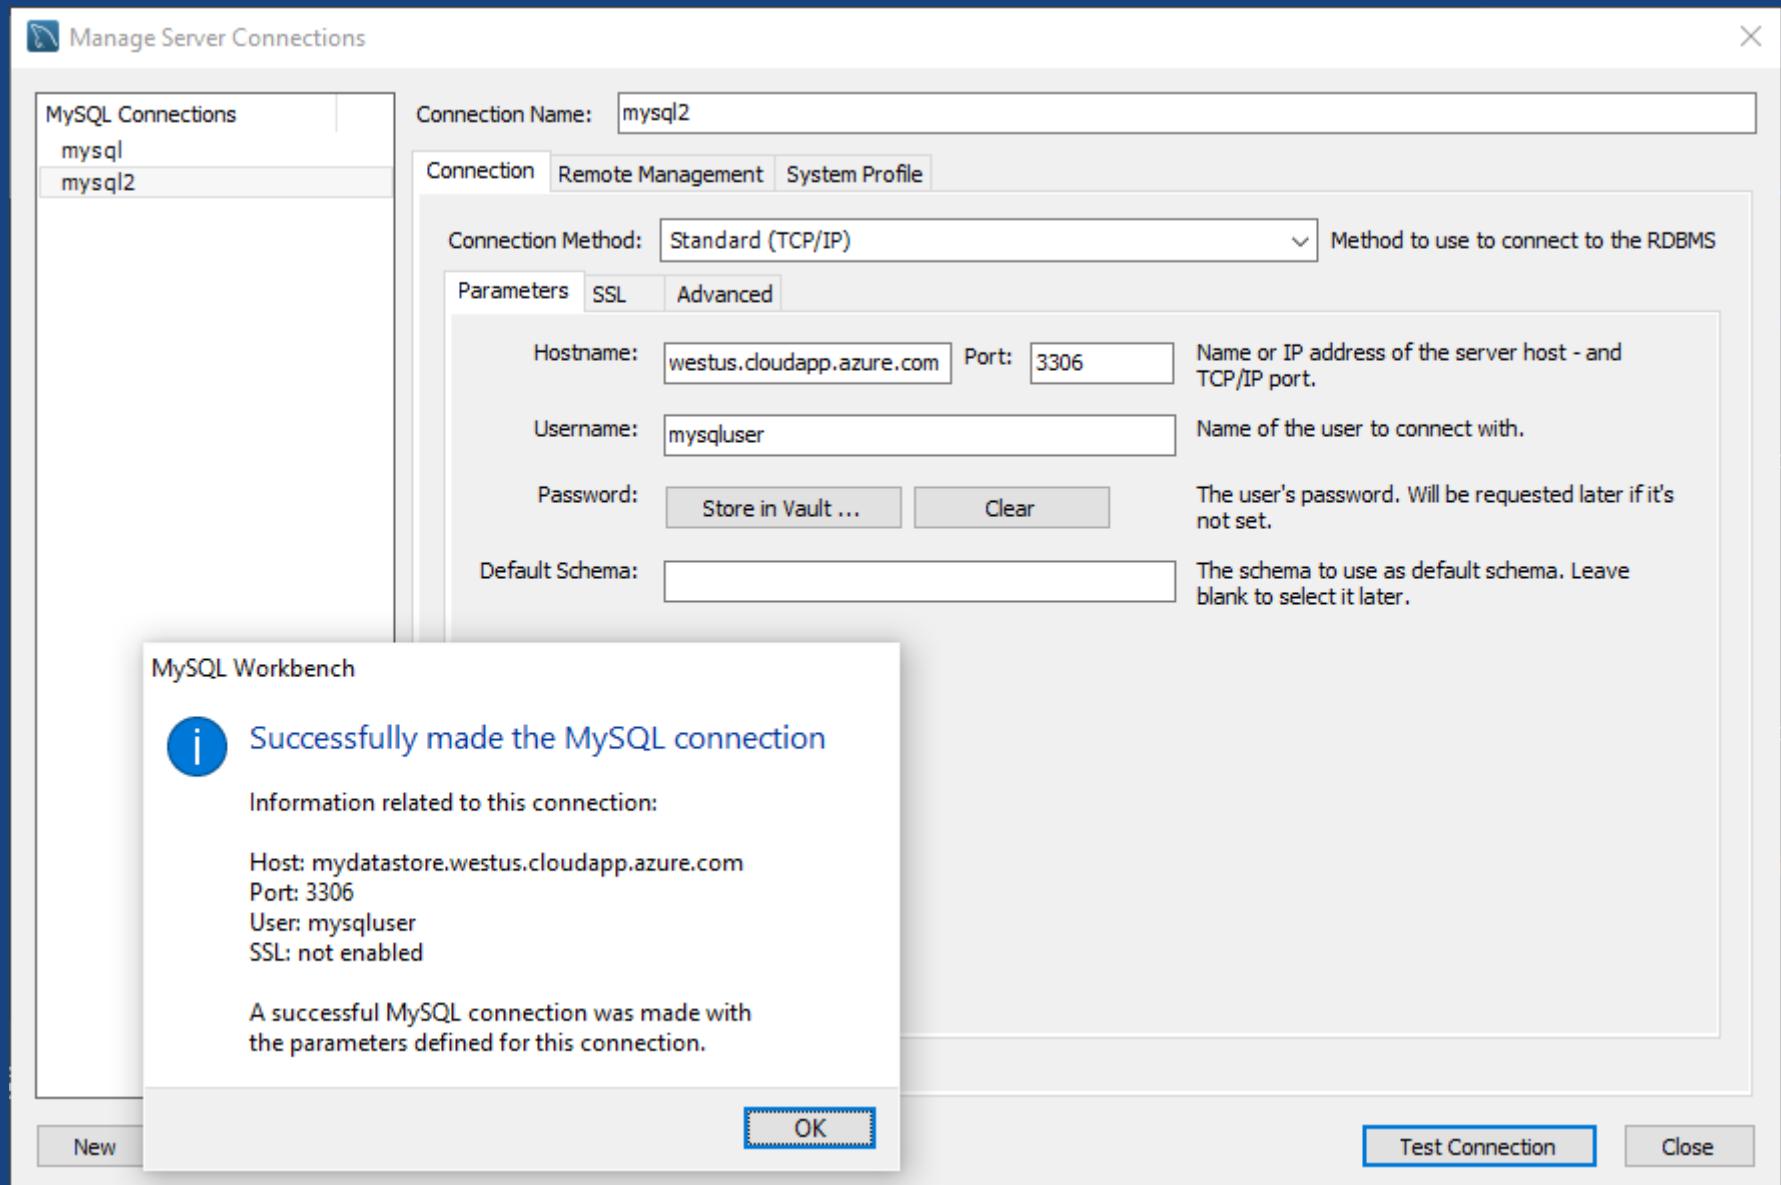
password (replace  
with a better one)

vim /etc/mysql/my.cnf

```
# Change the IP address  
# bind-address to 0.0.0.0  
  
# COMMENTED OUT  
# bind-address = 127.0.0.1  
  
bind-address = 0.0.0.0
```

reboot after  
edit

# Creating a user



# Creating a user

```
mysql> create database azurecoure;  
mysql> create user 'mysqluser'@'localhost' identified by 'password';  
mysql> grant all on azurecoure.* to 'mysqluser'@'%' identified by 'password';
```

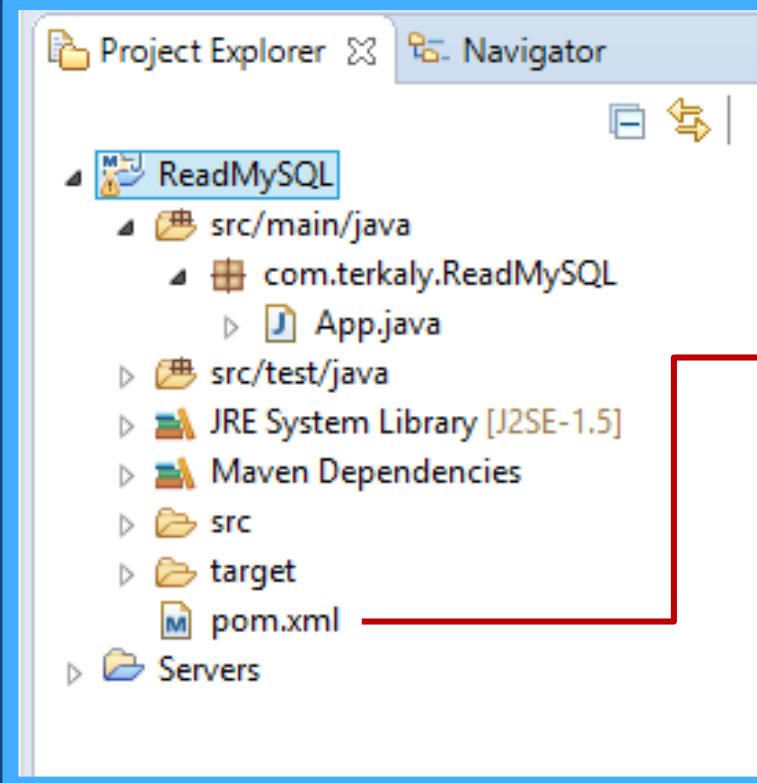
<http://www.yougetsignal.com/tools/open-ports/>

# Listing Open Ports

```
$ azure vm endpoint list oreillytestvm
```

Name of VM (you can  
see it on the command  
line)

# Setting Up MySQL/JDBC Driver on Java Client



The screenshot shows the Eclipse IDE's Project Explorer view. A red box highlights the 'pom.xml' file under the 'src' folder. A red arrow points from this highlighted file to the dependency code shown in the large black box.

```
<dependency>
    <groupId>mysql</groupId>
    <artifactId>mysql-connector-java</artifactId>
    <version>5.1.36</version>
</dependency>
```

# Inserting Data with Java into MySQL

```
package com.terkaly.ReadMySQL;

import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.PreparedStatement;
import java.sql.ResultSet;
import java.sql.SQLException;
import java.sql.Statement;

public session App {
    static private Connection connect = null;
    static private Statement statement = null;
    static private PreparedStatement preparedStatement = null;
    static private ResultSet resultSet = null;
    public static void main(String[] args) throws
        InstantiationException,
        IllegalAccessException, ClassNotFoundException {
        try {
            Class.forName("com.mysql.jdbc.Driver").newInstance();
            // Setup the connection with the DB
            connect = DriverManager.getConnection(
                "jdbc:mysql://104.42.126.213:3306/azuresession?" +
                "user=root&password=root&" +
                "useUnicode=true&characterEncoding=UTF-8");
        }
    }
}
```

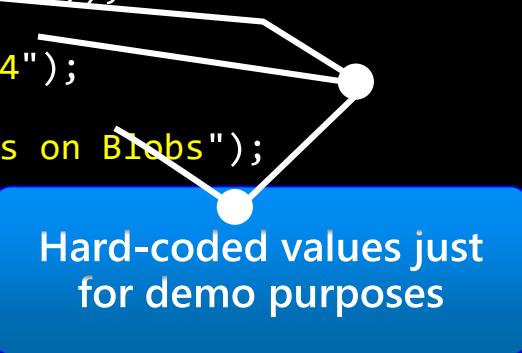
IP Address or DNS Name of  
your MySQL VM

Database = azuresession

Continued...

# Inserting Data with Java into MySQL

```
String query =  
    "insert into sessions(id, sessionnumber, sessiontitle)  
" +  
    "values(?, ?, ?);  
PreparedStatement preparedStmt =  
    connect.prepareStatement(query);  
preparedStmt.setInt(1, 4);  
preparedStmt.setString(2, "0404");  
preparedStmt.setString(3,  
    "Performing CRUD Operations on Blobs");  
preparedStmt.executeUpdate();  
connect.close();  
} catch (Exception e) {  
    e.printStackTrace();  
}  
}
```



Hard-coded values just  
for demo purposes

1. A prepared statement or parameterized statement is a feature used to execute the same or similar database statements repeatedly with high efficiency

2. It is used with SQL statements such as queries or updates

3. It takes the form of a template into which certain constant values are substituted during each execution

# Selecting data using Java to query MySQL

```
public static void PerformSelectOfData() throws SQLException {  
    // Statements allow to issue SQL queries to the database  
    Statement statement;  
    try {  
        statement = connect.createStatement();  
        resultSet =  
            statement.executeQuery("select * from azuresession.sessions");  
        WriteResults(resultSet);  
    } catch (Exception e) {  
        e.printStackTrace();  
    }  
    // Result set get the result of the SQL query  
}
```

# Displaying data using Java to query MySQL

```
public static void WriteResults(ResultSet resultSet) throws SQLException {
    System.out.println("The columns in the table are: ");

    try {
        System.out.println("Table: " + resultSet.getMetaData().getTableName(1));
        for (int i = 1; i <= resultSet.getMetaData().getColumnCount(); i++) {
            System.out.println("Column " + i + " " +
                               resultSet.getMetaData().getColumnName(i));
        }

        while (resultSet.next()) {
            int id = resultSet.getInt("id");
            String sessionnumber = resultSet.getString("sessionnumber");
            String sessiontitle = resultSet.getString("sessiontitle");
            System.out.println("ID: " + String.valueOf(id));
            System.out.println("Session Number: " + sessionnumber);
            System.out.println("Session Title: " + sessiontitle);
        }
    } catch (Exception e) {
        e.printStackTrace();
    }
}
```

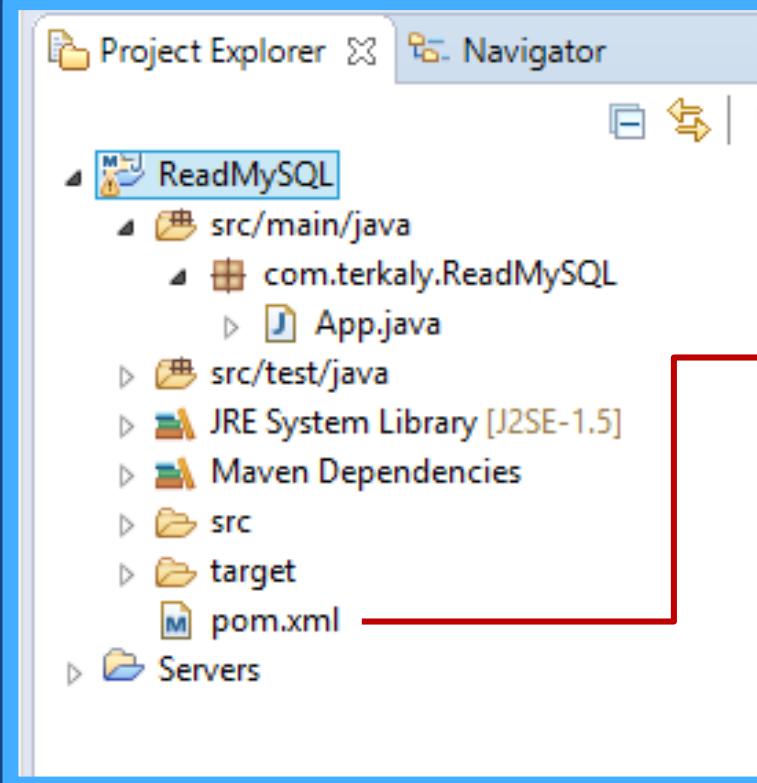
# **CONNECTING TO MYSQL FROM JAVA – PART 2**

**O'REILLY®**

# **CONNECTING TO MYSQL FROM JAVA – PART 2 (HANDS-ON)**

**O'REILLY®**

# Setting Up MySQL/JDBC Driver on Java Client



The screenshot shows the Eclipse IDE's Project Explorer view. A red box highlights the 'pom.xml' file under the 'src' folder. A red arrow points from this highlighted file to the dependency code shown in the large black box.

```
<dependency>
    <groupId>mysql</groupId>
    <artifactId>mysql-connector-java</artifactId>
    <version>5.1.36</version>
</dependency>
```

# Inserting Data with Java into MySQL

```
package com.terkaly.ReadMySQL;

import java.sql.Connection;
import java.sql.DriverManager;
import java.sql.PreparedStatement;
import java.sql.ResultSet;
import java.sql.SQLException;
import java.sql.Statement;

public session App {
    static private Connection connect = null;
    static private Statement statement = null;
    static private PreparedStatement preparedStatement = null;
    static private ResultSet resultSet = null;
    public static void main(String[] args) throws
        InstantiationException,
        IllegalAccessException, ClassNotFoundException {
        try {
            Class.forName("com.mysql.jdbc.Driver").newInstance();
            // Setup the connection with the DB
            connect = DriverManager.getConnection(
                "jdbc:mysql://104.42.126.213:3306/azuresession?" +
                "user=root&password=root&" +
                "useUnicode=true&characterEncoding=UTF-8");
        }
    }
}
```

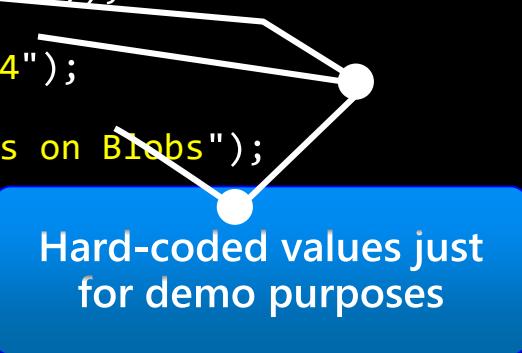
IP Address or DNS Name of  
your MySQL VM

Database = azuresession

Continued...

# Inserting Data with Java into MySQL

```
String query =  
    "insert into sessions(id, sessionnumber, sessiontitle)  
" +  
    "values(?, ?, ?);  
PreparedStatement preparedStmt =  
    connect.prepareStatement(query);  
preparedStmt.setInt(1, 4);  
preparedStmt.setString(2, "0404");  
preparedStmt.setString(3,  
    "Performing CRUD Operations on Blobs");  
preparedStmt.executeUpdate();  
connect.close();  
} catch (Exception e) {  
    e.printStackTrace();  
}  
}
```



Hard-coded values just  
for demo purposes

1. A prepared statement or parameterized statement is a feature used to execute the same or similar database statements repeatedly with high efficiency

2. It is used with SQL statements such as queries or updates

3. It takes the form of a template into which certain constant values are substituted during each execution

# Selecting data using Java to query MySQL

```
public static void PerformSelectOfData() throws SQLException {
    // Statements allow to issue SQL queries to the database
    Statement statement;
    try {
        statement = connect.createStatement();
        resultSet =
            statement.executeQuery("select * from azuresession.sessions");
        WriteResults(resultSet);
    } catch (Exception e) {
        e.printStackTrace();
    }
    // Result set get the result of the SQL query
}
```

# Displaying data using Java to query MySQL

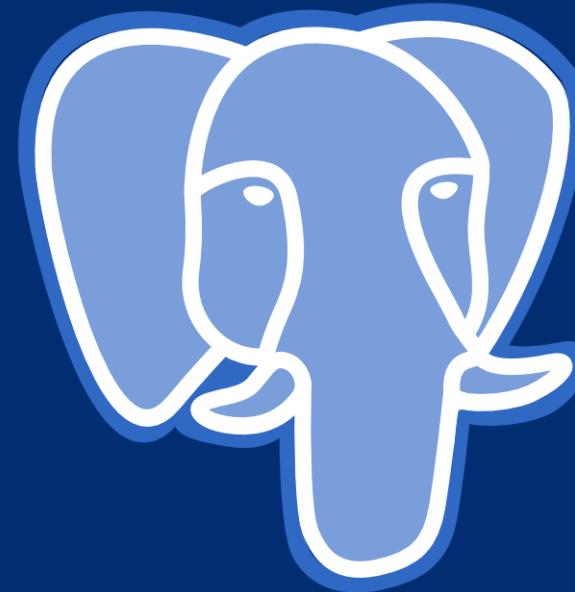
```
public static void WriteResults(ResultSet resultSet) throws SQLException {  
    System.out.println("The columns in the table are: ");  
  
    try {  
        System.out.println("Table: " + resultSet.getMetaData().getTableName(1));  
        for (int i = 1; i <= resultSet.getMetaData().getColumnCount(); i++) {  
            System.out.println("Column " + i + " " +  
                               resultSet.getMetaData().getColumnName(i));  
        }  
  
        while (resultSet.next()) {  
            int id = resultSet.getInt("id");  
            String sessionnumber = resultSet.getString("sessionnumber");  
            String sessiontitle = resultSet.getString("sessiontitle");  
            System.out.println("ID: " + String.valueOf(id));  
            System.out.println("Session Number: " + sessionnumber);  
            System.out.println("Session Title: " + sessiontitle);  
        }  
    } catch (Exception e) {  
        e.printStackTrace();  
    }  
}
```

# INTRODUCING POSTGRES

O'REILLY®

# Apache PostGres

- PostgreSQL is an object-relational database management system (RDBMS)
- PostgreSQL is cross-platform and runs on many operating systems including:
  - Linux
  - FreeBSD
  - OS X
  - Solaris
  - Microsoft Windows



## PostgreSQL

- Protects against dirty reads and full serializability
- Handles complex SQL queries

# Apache PostGres

- Supports indexing methods that are not available in other databases
  - Updateable views
  - Materialized views
  - Triggers
  - Foreign keys
- A view is a virtual table representing the result of a database query
  - An updatable view is useful because you can provide a stable interface to a client and still change the structure of underlying tables

## Materialized View

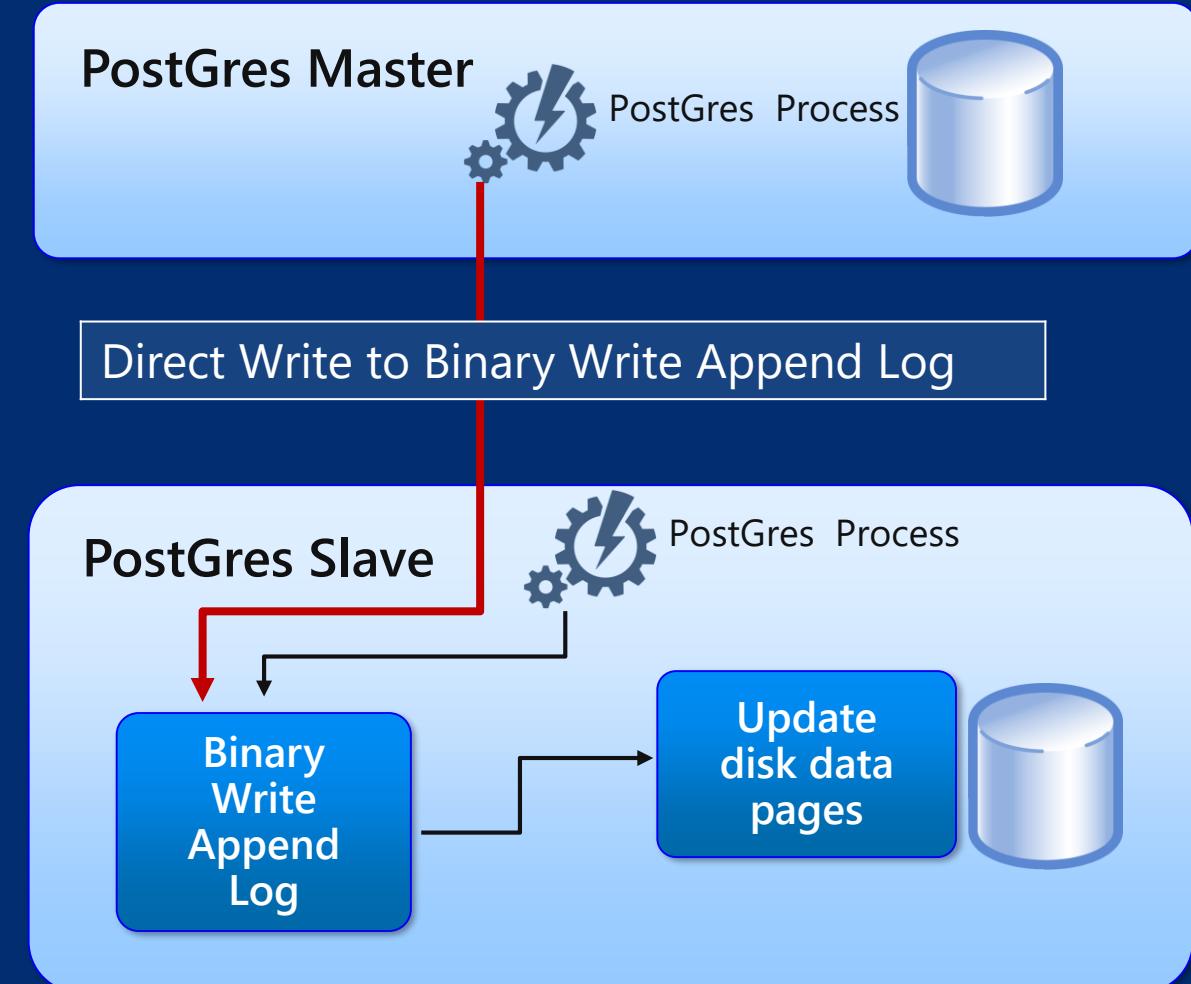
- A object that contains the results of a query
- Could be:
  - A local copy of data located remotely
  - A subset of the rows and/or columns of a table or join result
  - A summary using an aggregate function.

# Apache PostGres

## Built-in replication

- Supports a complete replica of your database
- Prevents all but a couple seconds of data loss
- Even under catastrophic circumstances
- Load balance between write master server and multiple read-only slave servers
- Run reporting or other long-running queries on a replica server
- Reduce load on main transaction-processing server

## Write-ahead logging



# Apache PostGres

## User Defined Types

- The ability to create types
- You specify an existing base
- You can create your own base types in C
- Useful when defining constraints on a type
- Example
  - CREATE DOMAIN custInitials AS `text` CHECK (value ~ '^[a-zA-Z][a-zA-Z]\$');

You can create your own base types beyond scalar types, like "text," "double" etc



Creating a new base type requires implementing functions to operate on the type in a low-level language, usually C.

Could also be enums

- CREATE TYPE shirtsize AS ENUM ('large', 'medium', 'small');
- You must define the input and output functions

# Apache PostGres

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

## Tunable Levels of Durability

- Transaction durability is tunable - can be specified per-database, per-user, per-session or even per-transaction
- Some workloads need to favor performance over data integrity
- There can be a mixture of synchronous and asynchronous standby servers

- Shared Buffers
- Max Connections
- listen\_addresses
- effective\_cache\_size
- checkpoint\_segments
- autovacuum
- logging
- wal\_sync\_method/wal\_buffers
- max\_prepared\_transactions
- synchronous\_commit
- random\_page\_cost

# Apache PostGres

How to tune and tweek PostGres  
Setting = Shared Buffers

- How much memory is dedicated to caching data
- Avoid need to recompile Linux kernel for large chunks of shared memory
- PostGres starts with 25% as a target for percentage of memory dedicated to shared buffers
- Sometimes OS-level caching might be more performant

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

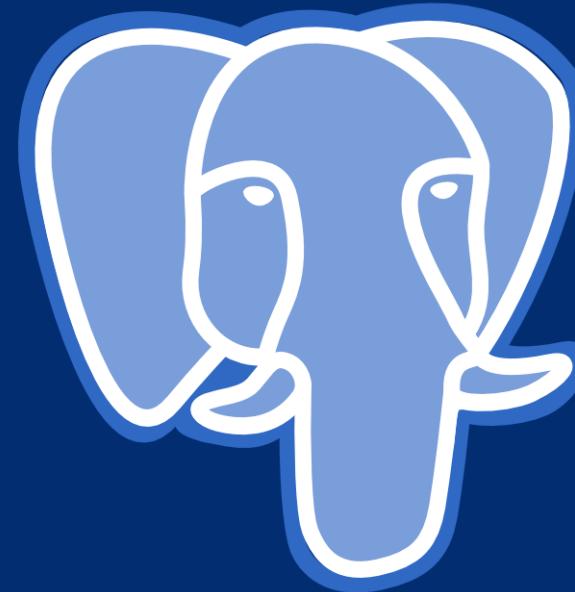
- Shared Buffers 
- Max Connections
- listen\_addresses
- effective\_cache\_size
- checkpoint\_segments
- autovacuum
- logging
- wal\_sync\_method/wal\_buffers
- max\_prepared\_transactions
- synchronous\_commit
- random\_page\_cost

# INTRODUCING POSTGRES

O'REILLY®

# Apache PostGres

- PostgreSQL is an object-relational database management system (RDBMS)
- PostgreSQL is cross-platform and runs on many operating systems including:
  - Linux
  - FreeBSD
  - OS X
  - Solaris
  - Microsoft Windows



## PostgreSQL

- Protects against dirty reads and full serializability
- Handles complex SQL queries

# Apache PostGres

- Supports indexing methods that are not available in other databases
  - Updateable views
  - Materialized views
  - Triggers
  - Foreign keys
- A view is a virtual table representing the result of a database query
  - An updatable view is useful because you can provide a stable interface to a client and still change the structure of underlying tables

## Materialized View

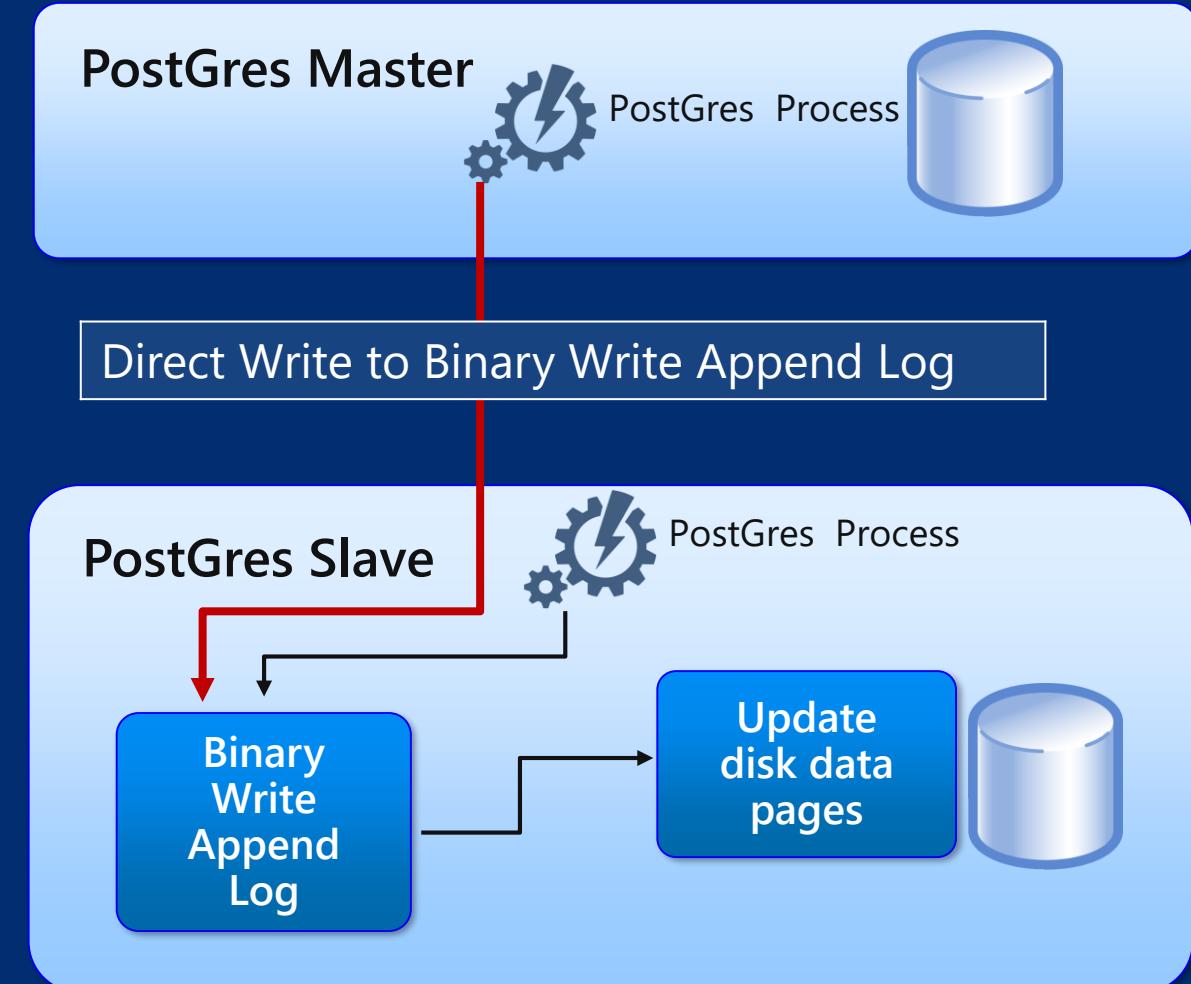
- A object that contains the results of a query
- Could be:
  - A local copy of data located remotely
  - A subset of the rows and/or columns of a table or join result
  - A summary using an aggregate function.

# Apache PostGres

## Built-in replication

- Supports a complete replica of your database
- Prevents all but a couple seconds of data loss
- Even under catastrophic circumstances
- Load balance between write master server and multiple read-only slave servers
- Run reporting or other long-running queries on a replica server
- Reduce load on main transaction-processing server

## Write-ahead logging



# Apache PostGres

## User Defined Types

- The ability to create types
- You specify an existing base
- You can create your own base types in C
- Useful when defining constraints on a type
- Example
  - CREATE DOMAIN custInitials AS `text` CHECK (value ~ '^[a-zA-Z][a-zA-Z]\$');

You can create your own base types beyond scalar types, like "text," "double" etc



Creating a new base type requires implementing functions to operate on the type in a low-level language, usually C.

Could also be enums

- CREATE TYPE shirtsize AS ENUM ('large', 'medium', 'small');
- You must define the input and output functions

# Apache PostGres

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

## Tunable Levels of Durability

- Transaction durability is tunable - can be specified per-database, per-user, per-session or even per-transaction
- Some workloads need to favor performance over data integrity
- There can be a mixture of synchronous and asynchronous standby servers

- Shared Buffers
- Max Connections
- listen\_addresses
- effective\_cache\_size
- checkpoint\_segments
- autovacuum
- logging
- wal\_sync\_method/wal\_buffers
- max\_prepared\_transactions
- synchronous\_commit
- random\_page\_cost

# Apache PostGres

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

How to tune and tweek PostGres  
Setting = Shared Buffers

- How much memory is dedicated to caching data
- Avoid need to recompile Linux kernel for large chunks of shared memory
- PostGres starts with 25% as a target for percentage of memory dedicated to shared buffers
- Sometimes OS-level caching might be more performant

- Shared Buffers 
- Max Connections
- listen\_addresses
- effective\_cache\_size
- checkpoint\_segments
- autovacuum
- logging
- wal\_sync\_method/wal\_buffers
- max\_prepared\_transactions
- synchronous\_commit
- random\_page\_cost

# INTRODUCING POSTGRES

## PART 2

O'REILLY®

# Apache PostGres

How to tune and tweek PostGres  
Setting = Max Connections

- Maximum number of client connections allowed
- On good hardware expect a few hundred connnections
- Consider connection pooling software to reduce the connection overhead

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

- Shared Buffers
- Max Connections ←
- listen\_addresses
- effective\_cache\_size
- checkpoint\_segments
- autovacuum
- logging
- wal\_sync\_method/wal\_buffers
- max\_prepared\_transactions
- synchronous\_commit
- random\_page\_cost

# Apache PostGres

How to tune and tweek PostGres  
Setting = `listen_addresses`

- By default, PostgreSQL only responds to connections from the local host.
- To support TCP/IP networking and connectivity outside the VM, you need to change `listen_addresses` from its default
- Change to `listen_addresses = '*'`
- And then control who can and cannot connect via the `pg_hba.conf` file

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

- Shared Buffers
- Max Connections
- `listen_addresses` ←
- `effective_cache_size`
- `checkpoint_segments`
- `autovacuum`
- `logging`
- `wal_sync_method/wal_buffers`
- `max_prepared_transactions`
- `synchronous_commit`
- `random_page_cost`

# Apache PostGres

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

How to tune and tweek PostGres  
Setting = `effective_cache_size`

- Memory for disk caching by the operating system and within the database itself
- Takes into account the OS and other applications
- Used by the Query Planner
- A Low effective cache size may lead to faulty index usage
- 1/2 of total memory would be a normal conservative setting
- 3/4 of memory is a more aggressive but still reasonable amount

- Shared Buffers
- Max Connections
- `listen_addresses`
- `effective_cache_size` ←
- `checkpoint_segments`
- `autovacuum`
- `logging`
- `wal_sync_method/wal_buffers`
- `max_prepared_transactions`
- `synchronous_commit`
- `random_page_cost`

# Apache PostGres

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

How to tune and tweek PostGres  
Setting = `checkpoint_segments`

- PostgreSQL writes new transactions to the database in files called WAL segments that are 16MB
- For normal systems set to 10
- For write heavy set to 32 (checkpoint every 512MB)
- For bulk loading can bump up to 64

- Shared Buffers
- Max Connections
- `listen_addresses`
- `effective_cache_size`
- `checkpoint_segments` ←
- autovacuum
- logging
- `wal_sync_method/wal_buffers`
- `max_prepared_transactions`
- `synchronous_commit`
- `random_page_cost`

# Apache PostGres

How to tune and tweek PostGres  
Setting = autovacuum

- Autovacuum process takes care of several maintenance chores inside your database that you really need
- You should set to true

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

- Shared Buffers
- Max Connections
- listen\_addresses
- effective\_cache\_size
- checkpoint\_segments
- autovacuum
- logging
- wal\_sync\_method/wal\_buffers
- max\_prepared\_transactions
- synchronous\_commit
- random\_page\_cost

# INTRODUCING POSTGRES

## PART 3

O'REILLY®

# Apache PostGres

How to tune and tweek PostGres  
Setting = logging

- pgFouine is a tool used to analyze postgresql logs
- PgCluu is an handy tool from the author of PgBadger, and is a PostgreSQL performances monitoring and auditing tool.
- log\_min\_error\_statement to see errors
- log\_min\_duration\_statement to find slow queries

and more

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

- Shared Buffers
- Max Connections
- listen\_addresses
- effective\_cache\_size
- checkpoint\_segments
- autovacuum
- logging 
- wal\_sync\_method/wal\_buffers
- max\_prepared\_transactions
- synchronous\_commit
- random\_page\_cost

# Apache PostGres

How to tune and tweek PostGres

Setting =  
wal\_sync\_method/wal\_buffers

- After every transaction, PostgreSQL forces a commit to disk out to its write-ahead log
- Increasing wal\_buffers from its tiny default of a small number of kilobytes is helpful for write-heavy systems
- Benchmarking generally suggests that just increasing to 1MB is enough for some large systems
- This data may become data as more powerful hardware is increasingly becoming available

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

- Shared Buffers
- Max Connections
- listen\_addresses
- effective\_cache\_size
- checkpoint\_segments
- autovacuum
- logging
- wal\_sync\_method/wal\_buffers ←
- max\_prepared\_transactions
- synchronous\_commit
- random\_page\_cost

# Apache PostGres

How to tune and tweek PostGres

Setting =  
max\_prepared\_transactions

- PREPARE TRANSACTION prepares the current transaction for two-phase commit
- Its purpose is to allow an external transaction manager to perform atomic global transactions across multiple databases
- PREPARE TRANSACTION is for distributed transactions across multiple servers, usually used by transaction monitors

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

- Shared Buffers
- Max Connections
- listen\_addresses
- effective\_cache\_size
- checkpoint\_segments
- autovacuum
- logging
- wal\_sync\_method/wal\_buffers
- max\_prepared\_transactions ←
- synchronous\_commit
- random\_page\_cost

# Apache PostGres

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

How to tune and tweek PostGres  
Setting = synchronous\_commit

- Specifies whether transaction commit will wait for WAL records to be written to disk before the command returns a "success" indication to the client
- Set synchronous\_commit to off when performance is more important than durability of a transaction.

- Shared Buffers
- Max Connections
- listen\_addresses
- effective\_cache\_size
- checkpoint\_segments
- autovacuum
- logging
- wal\_sync\_method/wal\_buffers
- max\_prepared\_transactions
- synchronous\_commit ←
- random\_page\_cost

# Apache PostGres

[https://wiki.postgresql.org/wiki/Tuning\\_Your\\_PostgreSQL\\_Server](https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server)

How to tune and tweek PostGres  
Setting = random\_page\_cost

- Meant to represent the relative cost of looking up one row (out of many) via sequential reads, vs. the cost of looking up a single row individually using random access (disk seeks)
  - It influences the planner's decisions to use:
    - table scan
    - composite indexes
    - simple indexes
- Shared Buffers
  - Max Connections
  - listen\_addresses
  - effective\_cache\_size
  - checkpoint\_segments
  - autovacuum
  - logging
  - wal\_sync\_method/wal\_buffers
  - max\_prepared\_transactions
  - synchronous\_commit
  - random\_page\_cost ←

# INSTALLING POSTGRES

O'REILLY®

# Apache PostGres



```
$ apt-get install postgresql  
$ apt-get install postgresql-contrib
```

- Using apt-get to perform the install

# Apache PostGres



```
$ su postgres  
$ createdb azuresession  
$ psql -d azuresession  
  
$ create table sessions(id int, sessionnumber varchar(128),  
sessiontitle varchar(512));
```

- Leverage postgres account
- Create database
- Get to sql prompt and create table inside azuresession database

# Apache PostGres



```
$ sudo -u postgres psql postgres  
$ \password postgres
```

- Set the password for the “postgres” account
- The “postgres” account is added during installation

We will need these credentials when we write our Java Client Code

# Apache PostGres



```
$ iptables -A INPUT -s 40.78.31.35 -p tcp --destination-port 5432 -m state --state NEW,ESTABLISHED -j ACCEPT
```

```
$ iptables -A OUTPUT -d 40.78.31.35 -p tcp --source-port 5432 -m state --state ESTABLISHED -j ACCEPT
```

- Open up ports
- May or may not be needed depending on how your VM is setup
- 40.78.31.35 is the public IP address of our VM
- PostGres uses the 5432 port for client connections

# Apache PostGres



```
$ vim /etc/postgresql/9.3/main/postgresql.conf  
$ vim /etc/postgresql/9.3/main/pg_hba.conf  
$ /etc/init.d/postgresql restart
```

```
55 #  
56  
57 # - Connection Settings -  
58  
59 listen_addresses = '*'          # what IP address(es) to listen on;
```

postgresql.conf

```
15 # Your host entry needs to look like you see it below  
16 host      all  all  all  trust
```

pg\_hba.conf

```
$ azure network nsg rule create -g oreilly-rg -a oreillybigdata -n cassandra-rule1  
-c Allow -p '*' -r Inbound -y 100 -f Internet -o '*' -e '*' -u '*'
```

- Adds a rule to network security

Resource Group	oreilly-rg
Network Security Group	oreillybigdata
Network security rule	Cassandra-rule1
What the rule does	Allows incoming/outgoing TCP/UDP traffic on VM

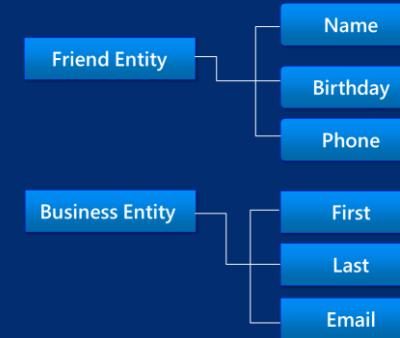
# NoSQL Data Stores

- Data storage and retrieval that is non-tabular/relational
- Benefits include simplicity, easier scaling
- Higher availability
- Key-value, graph, document

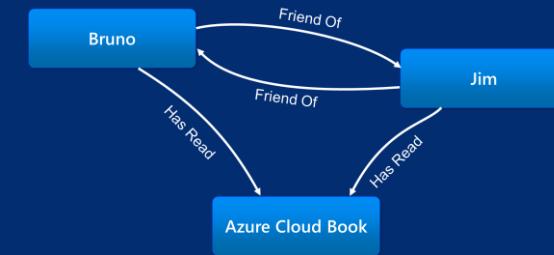
## Types of NoSQL

- Key-value

Demonstrated Previously



- Graph



- Document

See DocumentDB and MongoDB Hands-On

```
{  
  "data": [  
    {  
      "id": "X999_Y999",  
      "from": {  
        "name": "Bruno Terkaly",  
        "idsrc": "X12"  
      },  
      "to": {  
        "name": "Jim",  
        "idsrc": "X13"  
      }  
    }  
  ]  
}
```

# NoSQL Data Stores

- Not all data fits into 2 dimensions
- Volatility of the data model
- Likely to change and evolve?
- Schema rigidity
- Web-centric apps need flexibility

## Why or why not NoSQL

- Dependency on dedicated DBAs
- Developers like object focused nature of JSON
- Vertical scaling is expensive
- Distributed file systems
- Reporting needs
- Real time analytics

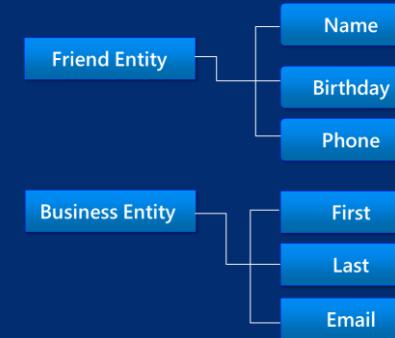
# NoSQL Data Stores

- Data storage and retrieval that is non-tabular/relational
- Benefits include simplicity, easier scaling
- Higher availability
- Key-value, graph, document

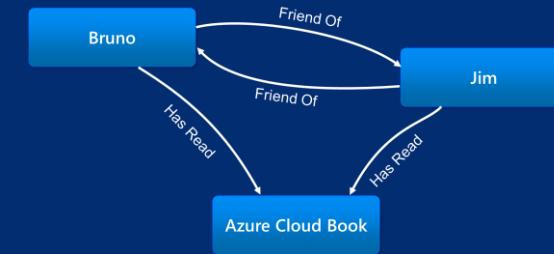
## Types of NoSQL

- Key-value

Demonstrated Previously



- Graph



- Document

See DocumentDB and MongoDB Hands-On

```
{  
  "data": [  
    {  
      "id": "X999_Y999",  
      "from": {  
        "name": "Bruno Terkaly",  
        "idsrc": "X12"  
      },  
      "to": {  
        "name": "Jim",  
        "idsrc": "X13"  
      }  
    }  
  ]  
}
```

# NoSQL Data Stores

- Not all data fits into 2 dimensions
- Volatility of the data model
- Likely to change and evolve?
- Schema rigidity
- Web-centric apps need flexibility

## Why or why not NoSQL

- Dependency on dedicated DBAs
- Developers like object focused nature of JSON
- Vertical scaling is expensive
- Distributed file systems
- Reporting needs
- Real time analytics

# Data Stores that run on Azure

- PaaS Offerings
  - DocumentDB (JSON)
  - Azure Tables (KeyValue)
- IaaS
  - Everything else – Azure is open
  - Cassandra, Mongo, etc

## Why DocumentDB?

- Need for managed service
- No pre-defined schema, or indexes
- A variety of clients
- Velocity and variability not known
- Scalable, fast deployment
- Low cost
- Query Language
- Support for transactions
- JSON-based
- REST-based access

# **PROVISIONING DOCUMENTDB DATABASES**

**O'REILLY®**

# Using the portal to interact with your data

- Add database
- Explore documents
- Execute queries
- View scripts
- Import data

## DocumentDB Portal

The screenshot shows the Azure DocumentDB portal interface. At the top, there's a header bar with the account name "azurecourse" and a "DocumentDB account". Below the header are several navigation icons: Settings, Add Database, Document Explorer, Query Explorer, Script Explorer, Import Data, Move, and Delete Account. The main area is divided into three sections:

- Databases:** Shows a "Document Explorer" sidebar with buttons for Create Document, Add Document, Refresh, and Settings. It lists a single database named "TestDB".
- Collections:** Shows a sidebar with a "Collections" dropdown set to "TestCollection".
- Documents:** Shows a "Documents" section with a search bar and a list of document IDs:
  - f053d9bd-8bde-4985-a11f-24585d8f3895
  - c7e28b0e-3c07-46a2-a1f-338beedb5e6b
  - 46faddb2-39b4-4684-add8-d25f4418a24b
  - 3ddeca52-7c6c-4134-b3e9-01f053bb83f0

# CONNECTING TO AZURE DOCUMENTDB WITH JAVA

O'REILLY®

# DocumentDB Code

- 1 Instantiate a DocumentClient w/ your DocumentDB Endpoint and AuthKey.
- 2 Instantiate a database object
- 3 Query for a database with the ID of TestDB
- 4 If does not exist, then create. Else get a reference to it.
- 5 Define a new collection using the id above.
- 6 Query for a collection called TestCollection
- 7 If collection not found, create
- 8 Create an object, serialize it in to JSON, and wrap it in to a document.
- 9 Insert document into document collection
- 10 Retrieve all inserted documents
- 11 If find at least one, retrieve the first one
- 12 Convert to a real object
- 13 Print out the first name

- DocumentClient documentClient
- Database myDatabase
- List<Database> databases
- DocumentCollection myCollection
- RequestOptions requestOptions
- SqlParameter sqlQuerySpec
- List<DocumentCollection> collections
- Person person = new Person();
- Document myDocument
- List<Document> documentList

# **CONNECTING TO AZURE DOCUMENTDB WITH JAVA – PART 2**

**O'REILLY®**

# DocumentDB Code

- 1 Instantiate a DocumentClient w/ your DocumentDB Endpoint and AuthKey.
- 2 Instantiate a database object
- 3 Query for a database with the ID of TestDB
- 4 If does not exist, then create. Else get a reference to it.
- 5 Define a new collection using the id above.
- 6 Query for a collection called TestCollection
- 7 If collection not found, create
- 8 Create an object, serialize it in to JSON, and wrap it in to a document.
- 9 Insert document into document collection
- 10 Retrieve all inserted documents
- 11 If find at least one, retrieve the first one
- 12 Convert to a real object
- 13 Print out the first name

- DocumentClient documentClient
- Database myDatabase
- List<Database> databases
- DocumentCollection myCollection
- RequestOptions requestOptions
- SqlParameter sqlQuerySpec
- List<DocumentCollection> collections
- Session session, Rating rating;
- Document myDocument
- List<Document> documentList

# **CONNECTING TO AZURE DOCUMENTDB WITH JAVA – PART 3**

**O'REILLY®**

# **CONNECTING TO AZURE DOCUMENTDB WITH JAVA – PART 4**

**O'REILLY®**

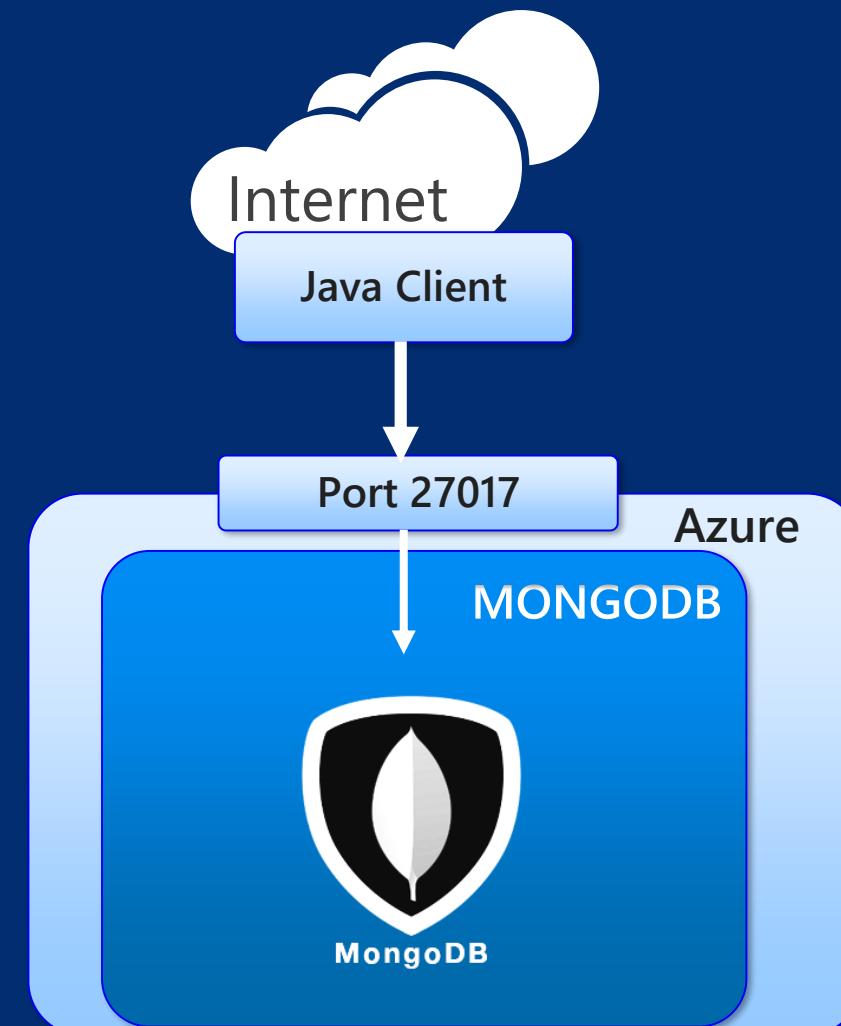
# **UNDERSTANDING MONGODB DATABASES**

**O'REILLY®**

# Understanding MongoDB

- MongoDB is a cross-platform document-oriented database
- Is a NoSQL database
- Data stored as JSON-like documents with dynamic schemas
- MongoDB calls the format BSON

## MongoDB From Java



# Understanding MongoDB

- A popular combination is Node.js + Mongo + JavaScript
- Makes the integration of data in certain types of applications easier and faster
- Used by Craigslist, eBay, and Foursquare among others
- As of July 2015, MongoDB is the fourth most popular type of database management system

## MongoDB and Node.js



- Node's is all about JavaScript
- This means that MongoDB documents get their most natural representation
  - JSON everywhere
    - the application layer
    - the database layer

# Understanding MongoDB

- Document-oriented
  - Can store relational objects in one document
- Supports Ad hoc queries
  - Supports search by field, range queries, regular expression searches
- Indexing
  - Any field in a MongoDB document can be indexed
- Replication
  - MongoDB provides high availability with replica sets
- Load balancing
  - MongoDB scales horizontally using sharding using a shard key
- Aggregation
  - MapReduce can be used for batch processing of data and aggregation operations
  - Similar to the SQL GROUP BY clause
- Server-side JavaScript execution
  - JavaScript can be used in queries and aggregation functions
- Capped collections
  - Acts as a circular queue on overflow

# PROVISIONING MONGODB DATABASES

O'REILLY®

# Installing Mongo

```
$ # Import the public key used by the package management system.  
$  
  
$ # Create a list file for MongoDB.  
$  
  
$ # Reload local package database.  
$  
  
$ # Install the latest stable version of MongoDB.  
$  
  
$ # Start MongoDB  
$
```

# Installing Mongo

```
$ # Verify port where MongoDB is listening  
$  
  
$ # Open up ports  
$  
  
$  
  
$ # Allow outside traffic in  
$  
$  
  
$ # Listen to local interface only. Comment out to listen on all interfaces.  
$ #COMMENT OUT -> bind_ip = 127.0.0.1  
$  
  
$  
$  
  
$ # Get the VMs Ip address from the portal  
$
```

The public IP  
Address of your VM

# **PROVISIONING MONGODB DATABASES**

## **PART 2**

**O'REILLY®**

# Installing Mongo

```
# BASH Provisioning Commands for MongoDB

```shell
# Import the public key used by the package management system.
sudo apt-key adv --keyserver hkp://keyserver.ubuntu.com:80 --recv 7F0CEB10

# Create a list file for MongoDB.
echo "deb http://repo.mongodb.org/apt/ubuntu $(lsb_release -sc)/mongodb-org/3.0 multiverse" | sudo tee /etc/apt/sources.list.d/mongodb-org-3.0.list

# Reload local package database.
sudo apt-get update

# Install the latest stable version of MongoDB.
sudo apt-get install -y mongodb-org

# Start MongoDB
sudo service mongod start

# Verify port where MongoDB is listening
vim /var/log/mongodb/mongod.log

# Open up ports
iptables -A INPUT -s 104.42.126.213 -p tcp --destination-port 27017 -m state --state NEW,ESTABLISHED -j ACCEPT
iptables -A OUTPUT -d 104.42.126.213 -p tcp --source-port 27017 -m state --state ESTABLISHED -j ACCEPT

// Allow outside traffic in
cd /etc
vim mongod.conf
# Listen to local interface only. Comment out to listen on all interfaces.
#bind_ip = 127.0.0.1
bind_ip = 0.0.0.0
rm /var/lib/mongodb/mongod.lock
reboot
#Should get your in
telnet 104.42.126.213 27017
```

# INTRODUCING APACHE CASSANDRA

O'REILLY®

# Apache Cassandra

- A free open source NoSQL database
- Designed to manage very large data sets
- Leverages distributed computing architectures using commodity hardware

Used by:

- CERN
- Comcast
- eBay
- GitHub
- GoDaddy
- Hulu
- Instagram
- Intuit
- Netflix
- Reddit
- The Weather Channel



# Apache Cassandra

Popular because:

- Almost linear scalability and high availability
- Massive ingestion capabilities
- Built in support for replication across data centers
- Hybrid between a key-value and a column-oriented (or tabular) database.
- Performance of log structured updates
- Built in caching



- Open source
- Peer to peer architecture
- Schema-Free
- High Availability
- Fault tolerant
- Tunable Consistency
- Column Oriented
- Elastic Scale

# Apache Cassandra

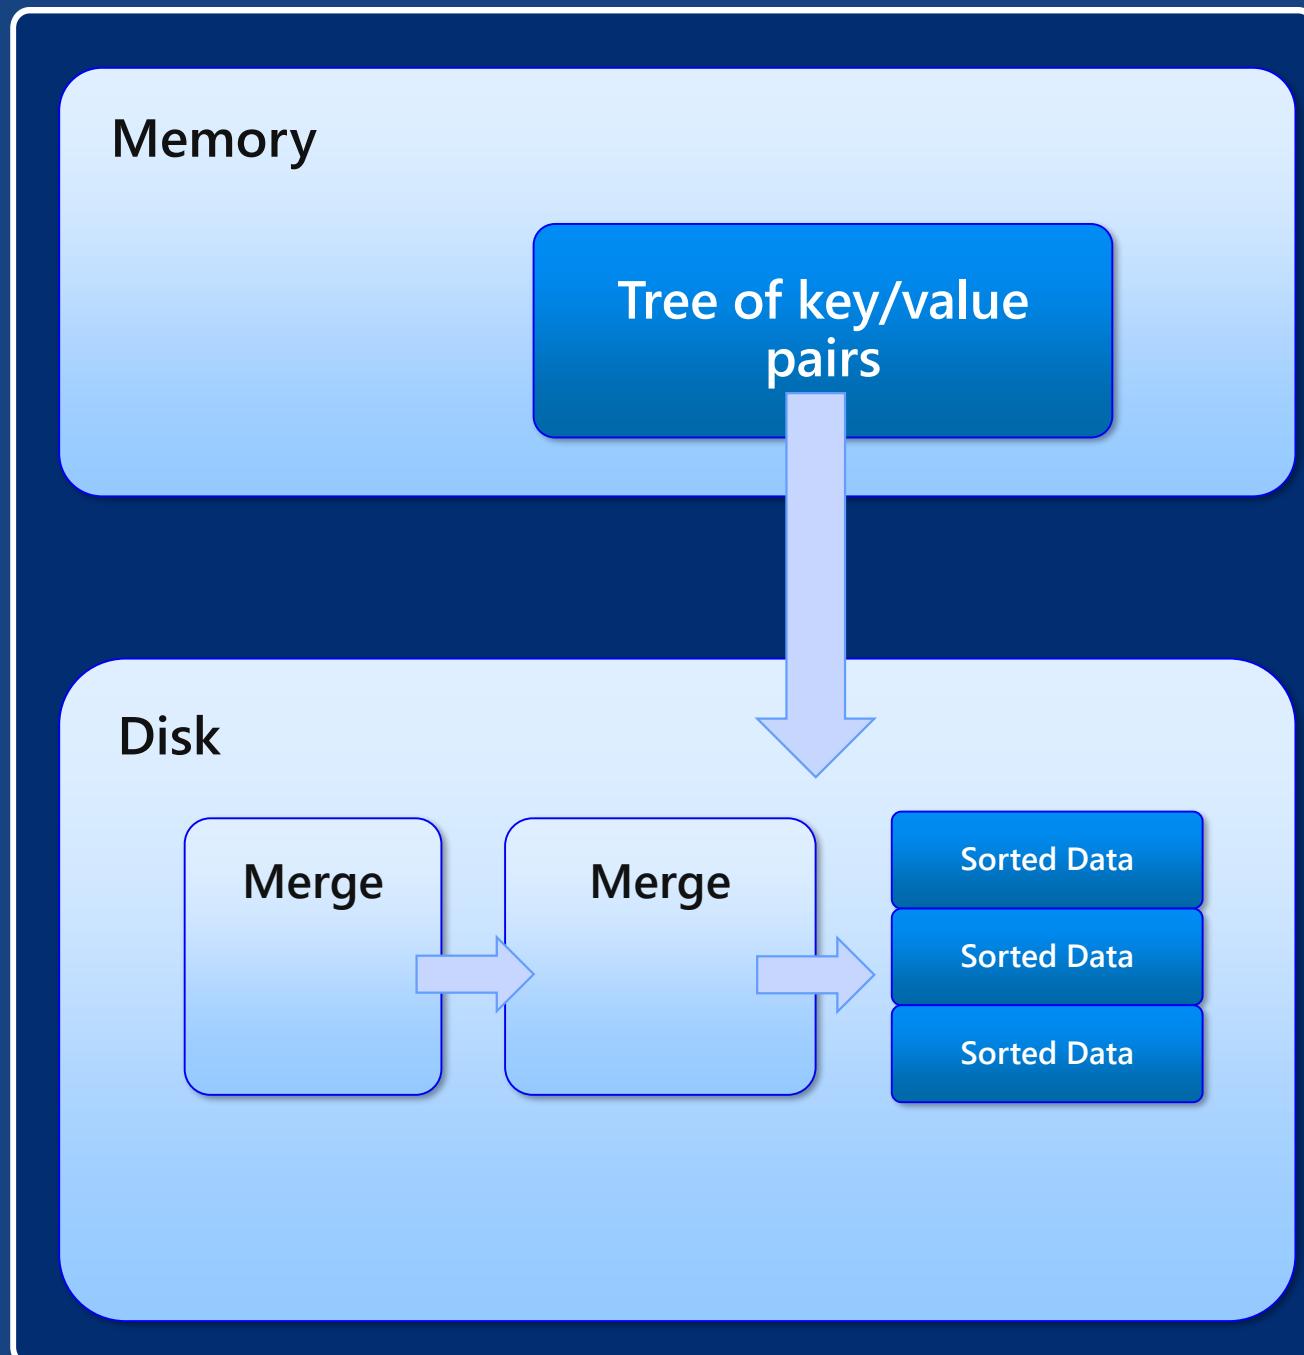
- A column-oriented DBMS
  - Stores data as columns
- Faster data retrieval because columns are closer together
- This reduces data read from disk, processed by the cpu, and cached in memory
- Advantages for data warehouses, CRM systems, etc
- Sweet spot is for ad hoc inquiry systems

- Imagine that I wish to analyze product sales by store
- Only a few columns are going to be used (location, product id, sales)
- Better that this data is grouped together for efficiency purposes



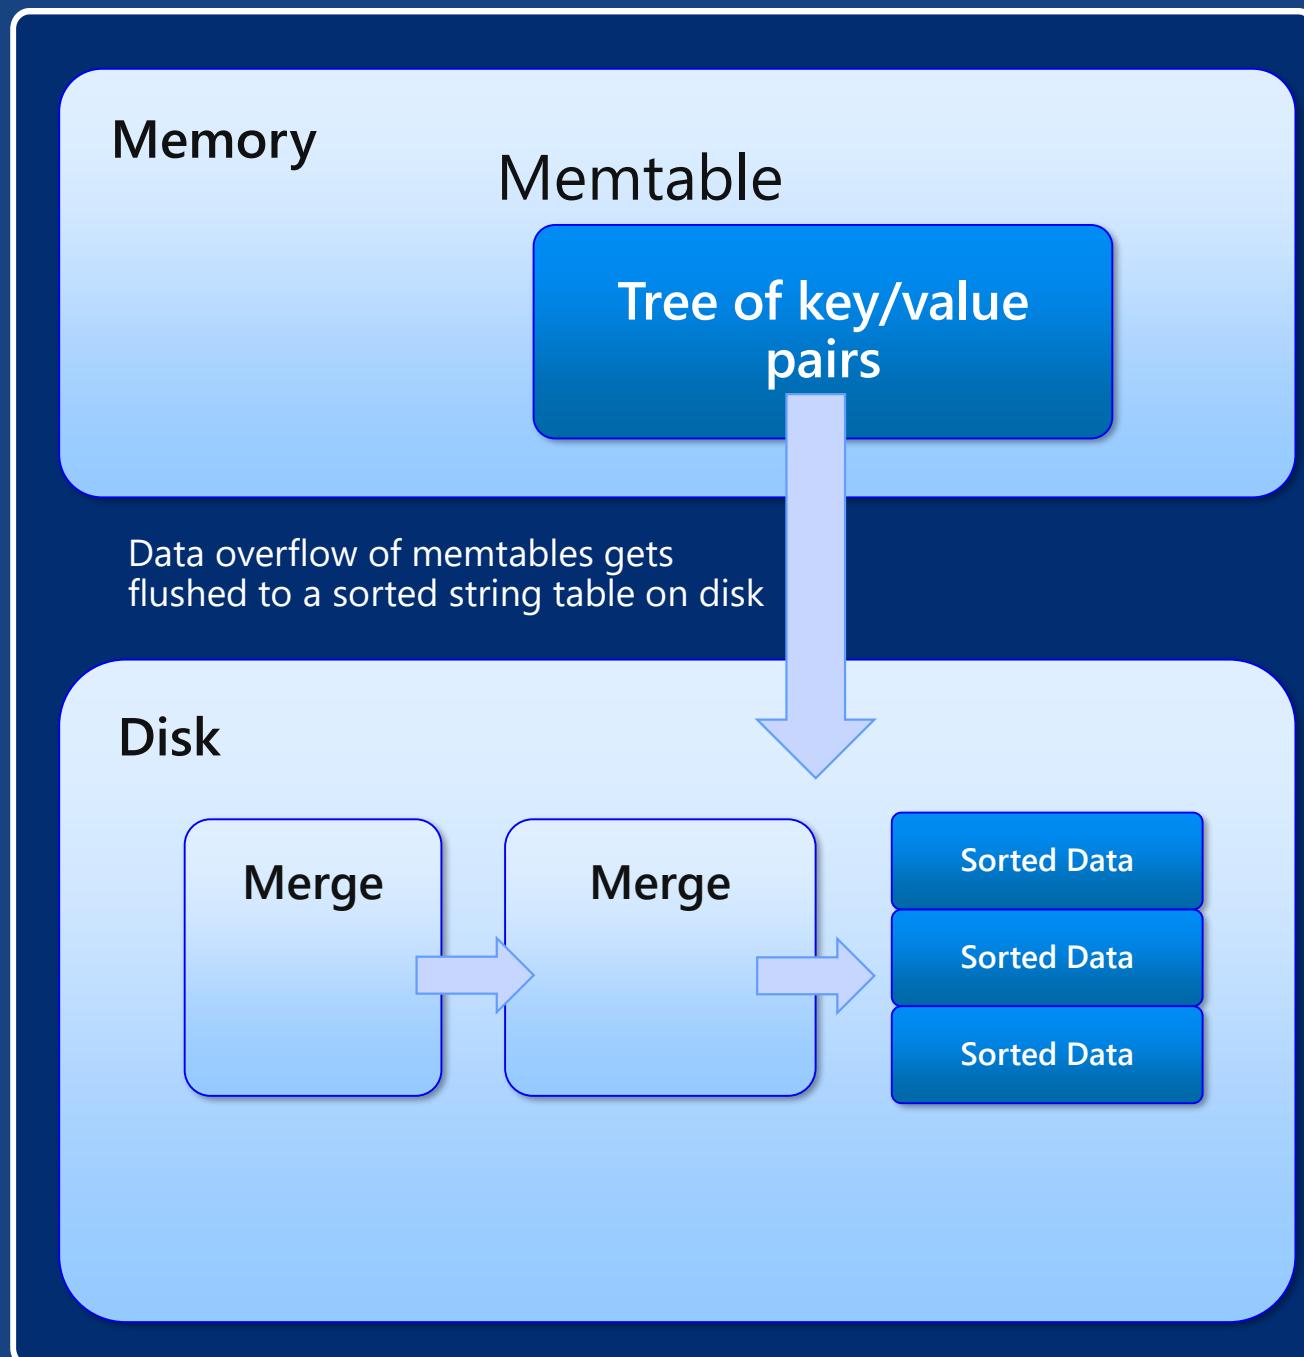
# Apache Cassandra

- Near linear scale ability for read/write operations
- No master nodes - all "created equal" nodes in cluster
- Consists of a log-structured merge tree
  - Standard disk-based index structures could be slow
  - The B-tree will dramatically increase the I/O cost of the transaction to maintain an index such in real time
  - It is very important to support a real-time index at high efficiency



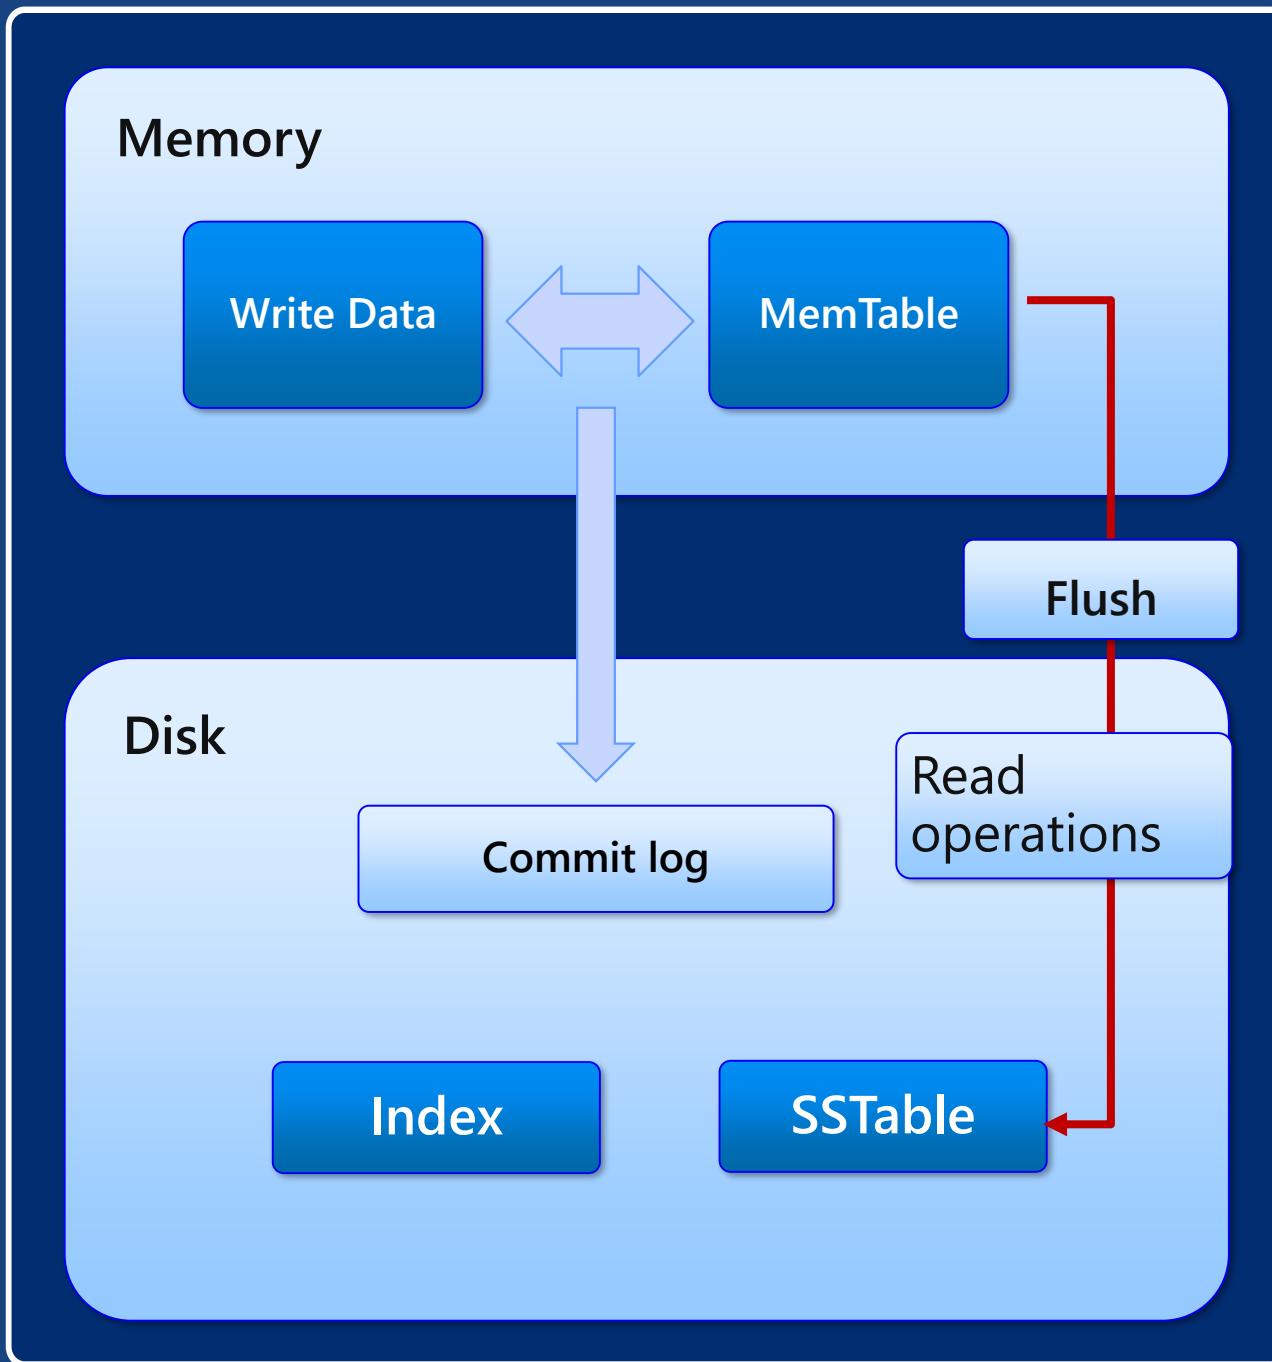
# Apache Cassandra

- Consists of a log-structured merge tree
  - A disk based data structure that defers and batches index changes
  - Reduces disk arm movement
  - Favors inserts over reads
  - Good for history of events, log files



# Apache Cassandra

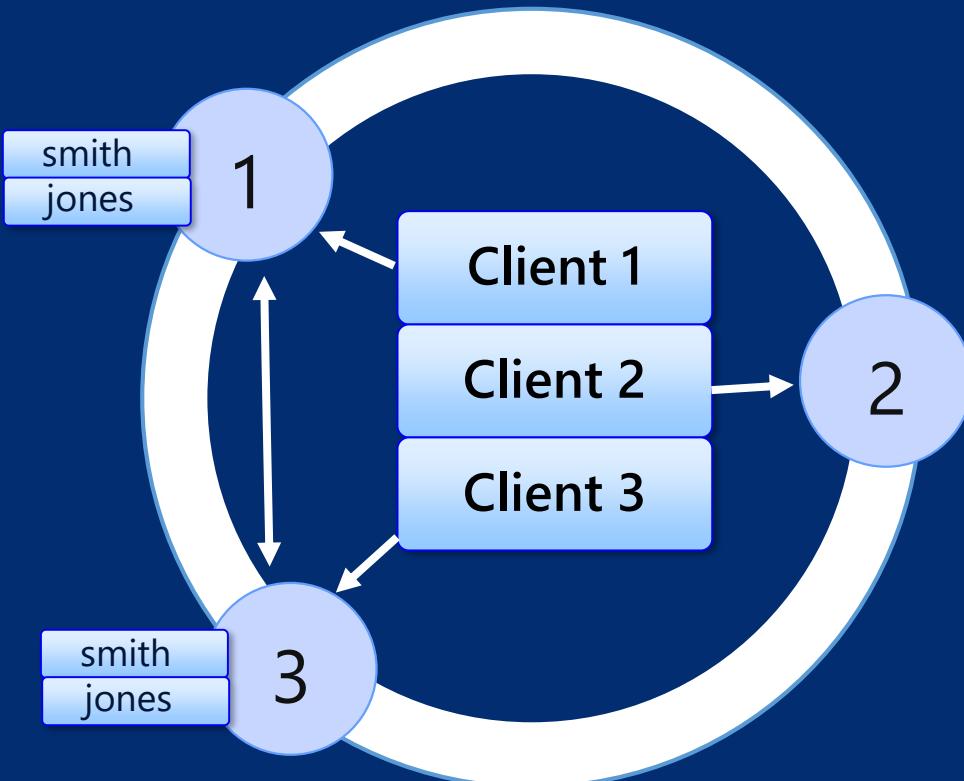
- Write operations are sent to a persistent commit log
- Also writes to a write-back cache called a memtable
- Data overflow of memtables gets flushed to a sorted string table on disk
- The sorted string tables get merged and compacted at specified intervals



# Apache Cassandra

- Append-only sequential writes for performance
- Data place randomly
- Periodically, the SSTable files are merged and compacted
- Ring network topology distributes data evenly
- Redundant writes to multiple nodes

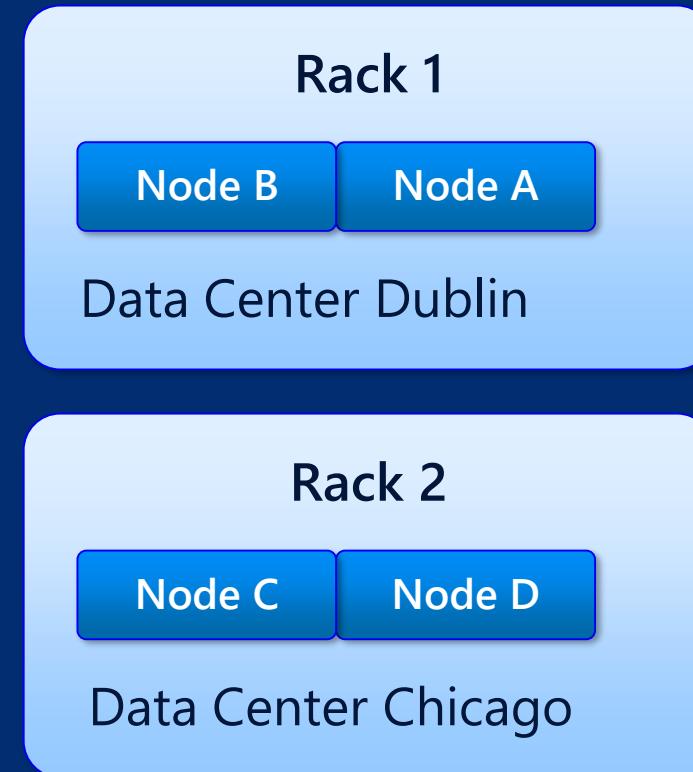
Partitioning Keys in the Cluster



# Apache Cassandra

- Has rack awareness for intelligent replica placement
- Support for automatic rebalance to avoid overutilization
- Node rebalance efficient because of limited data movement
- Allows you to think of your data as a range of database rows
- Automatically tracks unreachable replica nodes

## Cassandra Cluster



- RF-3 means every column stored in the database is stored 3 times.

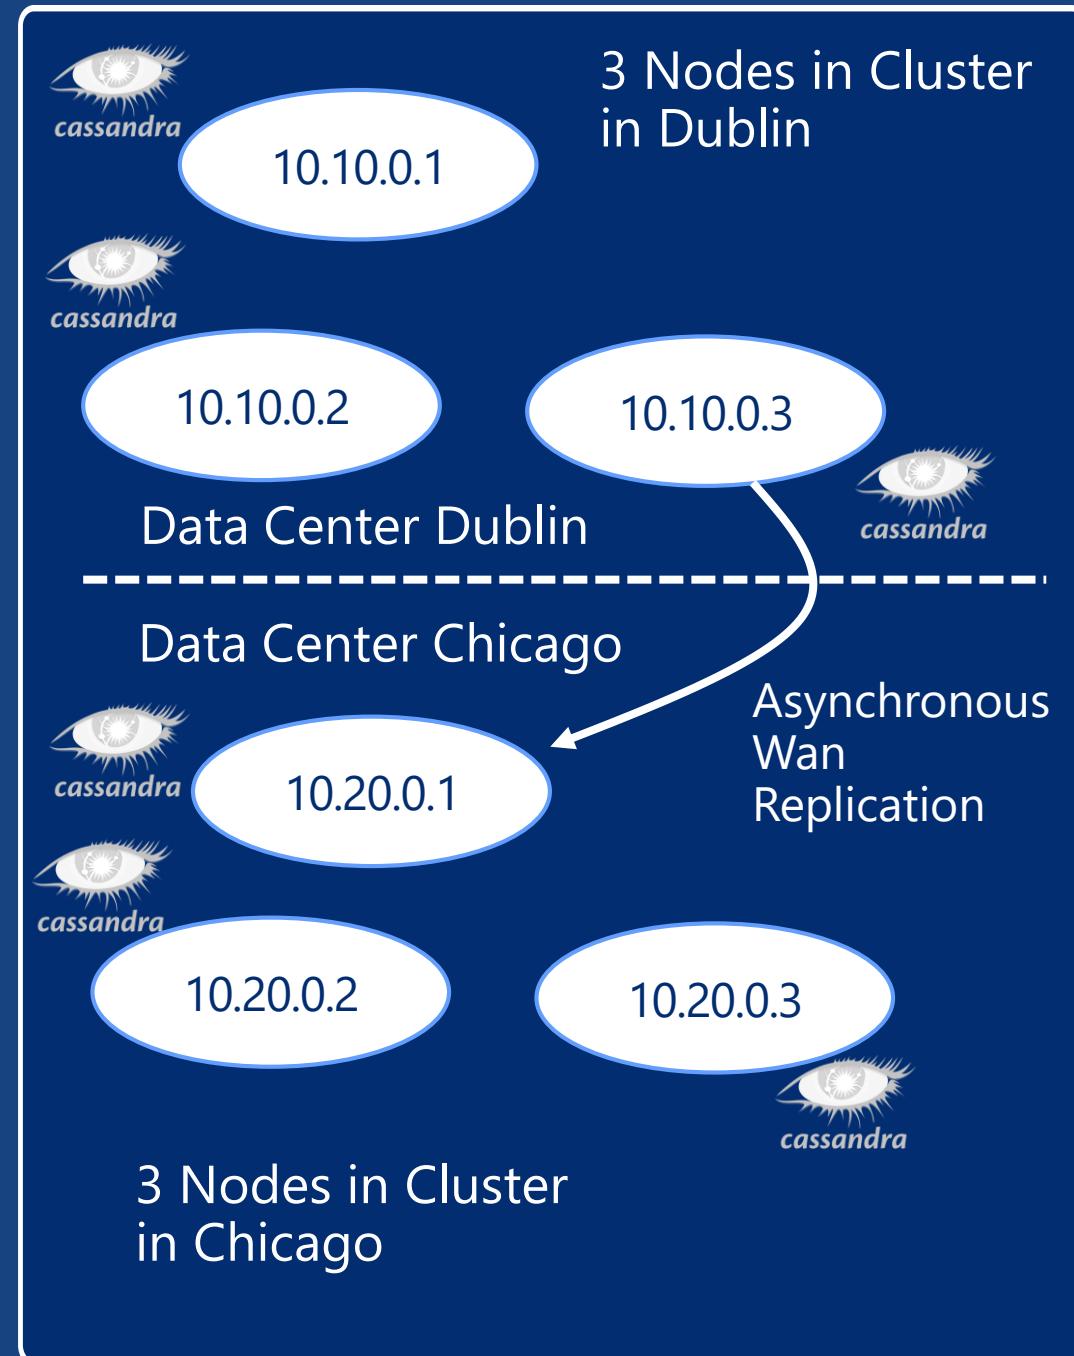
# Apache Cassandra

## Background Replication

- Issues background reads across replica nodes to ensure data consistency
- Will update replicas that are out of sync with each other

## Supports eventually consistent data model

- Offers an eventually consistent data model, sacrificing deterministic behavior for performance critical multiclient scenarios
- This means replicas are not updated simultaneously in a transactional manner



# Apache Cassandra

- Database engine **tunable** - level of consistency adjustable based on willingness to sacrifice read and write performance
- Confirmation of successful write operations can be delayed
- The most conservative setting is enforcing that all write operations to all replicas are complete before allowing clients to perform read operations
- A more balanced approach is to leverage the Quorum level of consistency whereby successful writes are acknowledgeable of a majority of replica nodes confirm write operations

## 11 Levels of consistency

- ALL
- EACH\_QUORUM
- QUORUM
- LOCAL\_QUORUM
- ONE
- TWO
- THREE
- LOCAL\_ONE
- ANY
- SERIAL
- LOCAL\_SERIAL

ALL	(1) A write must be written to the commit log and memtable on all replica nodes in the cluster for that partition. (2) Provides the highest consistency and the lowest availability of any other level.
ANY	(1) A write must be written to at least one node. Provides low latency and a guarantee that a write never fails. (2) Delivers the lowest consistency and highest availability.

# Apache Cassandra

- Supports a powerful query language (CQL) that includes SELECT, INSERT, UPDATE, and DELETE statements
- But no support for join operation and the referential integrity of traditional DBMS systems
- Supports common data types like integers, floats and doubles, blobs, and more
- Allows you to assign heavier work loads to nodes that are running on more powerful hardware

## Version 1.0 CQL Keywords

USE	DROP
SELECT	BATCH
UPDATE	CREATE KEYSPACE
DELETE	CREATE COLUMNFAMILY
TRUNCATE	CREATE INDEX

# Connecting to Cassandra

## Cassandra Ports

- 7000 for cluster communication
- 7001 if SSL is enabled
- 9160 for Thrift clients
- 9042 for native protocol clients
- 7199 for JMX



- The **Thrift** interface is a legacy API for older clients.
- JMX
  - Java Management Extensions
  - For managing and monitoring Java applications and services

# INSTALLING CASSANDRA

O'REILLY®

# Apache Cassandra



```
$ cd /tmp  
$ wget http://apache.mesi.com.ar/cassandra/2.1.9/apache-cassandra-2.1.9-bin.tar.gz  
$ mkdir -p ~/cassandra  
$ cd ~/cassandra  
$ tar -xvf /tmp/apache-cassandra-2.1.9-bin.tar.gz
```

- Download Cassandra with wget
- Untar contents

**Includes (1) the core server  
(2) the nodetool administration  
command-line interface  
(3) a development shell (cqlsh  
and the old cassandra-cli)**

# Apache Cassandra



```
$ vim ~/cassandra/apache-cassandra-2.1.9/conf/cassandra.yaml  
$ vim ~/.bashrc
```

```
446 rpc_address: 0.0.0.0  
360 # TCP port, for commands and data  
361 # For security reasons, you should not  
362 # expose this port to the internet.  
363 # Firewall it if needed.  
364 # used to be 7000  
365 storage_port: 8000  
366 ssl_storage_port: 8001
```

## cassandra.yaml

- Opening connectivity with `rpc_address`
- `storage_port` needed to be changed because of conflict

```
119 #  
120 # fix up environment variables and path  
121 #  
122 export M2_HOME="/usr/bin/apache-maven-3.3.3"  
123 export M2=$M2_HOME/bin  
124  
125 export CASSANDRA_HOME=~/cassandra/apache-cassandra-2.1.9  
126 export PATH="$M2:$CASSANDRA_HOME/bin:$PATH"
```

## .bashrc

- Setting `CASSANDRA_HOME` and the `PATH`

# Apache Cassandra



```
$ azure network nsg rule create -g oreilly-rg -a oreillybigdata -n cassandra-rule1  
-c Allow -p '*' -r Inbound -y 100 -f Internet -o '*' -e '*' -u '*'
```

- Adds a rule to network security

Resource Group	oreilly-rg
Network Security Group	oreillybigdata
Network security rule	Cassandra-rule1
What the rule does	Allows incoming/outgoing TCP/UDP traffic on VM

# Apache Cassandra



```
<dependency>
    <groupId>com.datastax.cassandra</groupId>
    <artifactId>cassandra-driver-core</artifactId>
    <version>2.1.8</version>
</dependency>
```

- The POM.xml entry for Cassandra
- <http://mavenrepository.com/artifact/com.datastax.cassandra/cassandra-driver-core/2.1.8>

# Apache Cassandra



```
// Adding the import statement  
  
import com.datastax.driver.core.*;  
public static void main( String[] args )  
{  
    try {  
        // Connect to Cassandra  
        connect("40.78.31.35");  
        // Create keyspace and table  
        createSchema();  
        // Load table with data  
        loadData();  
        // Query table using "where" clause  
        querySchema();  
        // cleanup and close  
        close();  
    }  
    catch (Exception e) {  
        e.printStackTrace();  
    }  
}
```

# Apache Cassandra



```
// connect to connect to Cassandra running in Azure
public static void connect(String node) {
    cluster = Cluster.builder().addContactPoint(node).build();
    Metadata metadata = cluster.getMetadata();
    System.out.printf("Connected to cluster: %s\n",
                      metadata.getClusterName());
    for ( Host host : metadata.getAllHosts() ) {
        System.out.printf("Datacenter: %s; Host: %s; Rack: %s\n",
                          host.getDatacenter(), host.getAddress(), host.getRack());
    }
    session = cluster.connect();
}
```

# Apache Cassandra



```
// Create the keyspace and tables
public static void createSchema() {
    session.execute("CREATE KEYSPACE IF NOT EXISTS simplex WITH replication " +
        "= {'session':'SimpleStrategy', 'replication_factor':3};");
    session.execute(
        "CREATE TABLE IF NOT EXISTS simplex.songs (" +
        "id uuid PRIMARY KEY," +
        "title text," +
        "album text," +
        "artist text," +
        "tags set<text>," +
        "data blob" +
        ");");
    ...
}
```

# Apache Cassandra



```
// load the data
public static void loadData() {

    session.execute(
        "INSERT INTO simplex.songs (id, title, album, artist, tags)  +
        "VALUES (" +
        "756716f7-2e54-4715-9f00-91dcbea6cf50," +
        "'La Petite Tonkinoise'," +
        "'Bye Bye Blackbird'," +
        "'Joséphine Baker'," +
        "{'jazz', '2013'})" +
        ");");

...
}
```

# Apache Cassandra



```
// query the data
public static void querySchema() {
    ResultSet results = session.execute("SELECT * FROM simplex.playlists "
        +
        "WHERE id = 2cc9ccb7-6221-4ccb-8387-f22b6a1b354d;");
    System.out.println(String.format("%-30s\t%-20s\t%-20s\n%s",
        "title",
        "album", "artist",
        "-----+-----"
        + "-----"));
    for (Row row : results) {
        System.out.println(String.format("%-30s\t%-20s\t%-20s",
            row.getString("title"),
            row.getString("album"), row.getString("artist")));
    }
    System.out.println();
}
```